

```

1 # Install necessary libraries
2 !pip install shap lime scikit-learn imbalanced-learn tensorflow pandas pyarrow polars
3
4 import os
5 import json
6 import numpy as np
7 import pandas as pd
8 import matplotlib.pyplot as plt
9 import seaborn as sns
10 from sklearn.model_selection import train_test_split
11 from sklearn.preprocessing import StandardScaler, OneHotEncoder
12 from sklearn.compose import ColumnTransformer
13 from sklearn.metrics import classification_report, roc_auc_score, precision_recall_fscore_support, roc_curve
14 from sklearn.ensemble import IsolationForest
15 from sklearn.inspection import permutation_importance
16 import shap
17 import tensorflow as tf
18 from tensorflow import keras
19

```

Requirement already satisfied: shap in /usr/local/lib/python3.12/dist-packages (0.48.0)
Requirement already satisfied: lime in /usr/local/lib/python3.12/dist-packages (0.2.0.1)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.12/dist-packages (1.6.1)
Requirement already satisfied: imbalanced-learn in /usr/local/lib/python3.12/dist-packages (0.14.0)
Requirement already satisfied: tensorflow in /usr/local/lib/python3.12/dist-packages (2.19.0)
Requirement already satisfied: pandas in /usr/local/lib/python3.12/dist-packages (2.2.2)
Requirement already satisfied: pyarrow in /usr/local/lib/python3.12/dist-packages (18.1.0)
Requirement already satisfied: polars in /usr/local/lib/python3.12/dist-packages (1.25.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (from shap) (2.0.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.12/dist-packages (from shap) (1.16.1)
Requirement already satisfied: tqdm>=4.27.0 in /usr/local/lib/python3.12/dist-packages (from shap) (4.67.1)
Requirement already satisfied: packaging>20.9 in /usr/local/lib/python3.12/dist-packages (from shap) (25.0)
Requirement already satisfied: slicer==0.0.8 in /usr/local/lib/python3.12/dist-packages (from shap) (0.0.8)
Requirement already satisfied: numba>=0.54 in /usr/local/lib/python3.12/dist-packages (from shap) (0.60.0)
Requirement already satisfied: cloudpickle in /usr/local/lib/python3.12/dist-packages (from shap) (3.1.1)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.12/dist-packages (from shap) (4.15.0)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages (from lime) (3.10.0)
Requirement already satisfied: scikit-image>=0.12 in /usr/local/lib/python3.12/dist-packages (from lime) (0.25.2)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (1.5.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (3.6.0)
Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.12/dist-packages (from tensorflow) (1.4.0)
Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.12/dist-packages (from tensorflow) (1.6.3)
Requirement already satisfied: flatbuffers>=24.3.25 in /usr/local/lib/python3.12/dist-packages (from tensorflow) (25.2.10)
Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>0.2.1 in /usr/local/lib/python3.12/dist-packages (from tensorflow) (0.2.0)
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.12/dist-packages (from tensorflow) (0.2.0)
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.12/dist-packages (from tensorflow) (18.1.1)
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.12/dist-packages (from tensorflow) (3.4.0)
Requirement already satisfied: protobuf!=4.21.0,!=4.21.1,!=4.21.2,!=4.21.3,!=4.21.4,!=4.21.5,<6.0.0dev,>=3.20.3 in /usr/local/lib/python3.12/dist-packages (from tensorflow) (4.21.6)
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.12/dist-packages (from tensorflow) (2.32.4)
Requirement already satisfied: setuptools in /usr/local/lib/python3.12/dist-packages (from tensorflow) (75.2.0)
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.12/dist-packages (from tensorflow) (1.17.0)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.12/dist-packages (from tensorflow) (3.1.0)
Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.12/dist-packages (from tensorflow) (1.17.3)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.12/dist-packages (from tensorflow) (1.74.0)
Requirement already satisfied: tensorboard>=2.19.0 in /usr/local/lib/python3.12/dist-packages (from tensorflow) (2.19.0)
Requirement already satisfied: keras>=3.5.0 in /usr/local/lib/python3.12/dist-packages (from tensorflow) (3.10.0)
Requirement already satisfied: h5py>=3.11.0 in /usr/local/lib/python3.12/dist-packages (from tensorflow) (3.14.0)
Requirement already satisfied: ml-dtypes<1.0.0,>=0.5.1 in /usr/local/lib/python3.12/dist-packages (from tensorflow) (0.5.3)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.12/dist-packages (from astunparse>=1.6.0->tensorflow) (0.37.1)
Requirement already satisfied: rich in /usr/local/lib/python3.12/dist-packages (from keras>=3.5.0->tensorflow) (13.9.4)
Requirement already satisfied: namex in /usr/local/lib/python3.12/dist-packages (from keras>=3.5.0->tensorflow) (0.1.0)
Requirement already satisfied: optree in /usr/local/lib/python3.12/dist-packages (from keras>=3.5.0->tensorflow) (0.17.0)
Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in /usr/local/lib/python3.12/dist-packages (from numba>=0.54->shap) (0.43.0)
Requirement already satisfied: charset_normalizer<4,>=2 in /usr/local/lib/python3.12/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.1.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.12/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.2.0)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.12/dist-packages (from requests<3,>=2.21.0->tensorflow) (2.0.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.12/dist-packages (from requests<3,>=2.21.0->tensorflow) (2025.2)
Requirement already satisfied: networkx>=3.0 in /usr/local/lib/python3.12/dist-packages (from scikit-image>=0.12->lime) (3.5)
Requirement already satisfied: pillow>=10.1 in /usr/local/lib/python3.12/dist-packages (from scikit-image>=0.12->lime) (11.3.0)
Requirement already satisfied: imageio!=2.35.0,>=2.33 in /usr/local/lib/python3.12/dist-packages (from scikit-image>=0.12->lime) (2.33.0)
Requirement already satisfied: tifffile>=2022.8.12 in /usr/local/lib/python3.12/dist-packages (from scikit-image>=0.12->lime) (2022.8.12)
Requirement already satisfied: lazy-loader>=0.4 in /usr/local/lib/python3.12/dist-packages (from scikit-image>=0.12->lime) (0.4)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.12/dist-packages (from tensorboard~>=2.19.0->tensorflow) (2.6.8)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.12/dist-packages (from tensorboard) (0.7.0)
Requirement already satisfied: werkzeug<-1.0.1 in /usr/local/lib/python3.12/dist-packages (from tensorboard~>=2.19.0->tensorflow) (0.12.2)

>Loading the CIC-IoT 2023 Dataset

```

1 # Load the CIC-IoT 2023 dataset
2 DATA_PATH_IOT_1 = "/content/drive/MyDrive/Datasets/CIC-IoT 2023/cic-iot_2023_1.csv"
3 DATA_PATH_IOT_2 = "/content/drive/MyDrive/Datasets/CIC-IoT 2023/cic-iot_2023_2.csv"
4 DATA_PATH_IOT_3 = "/content/drive/MyDrive/Datasets/CIC-IoT 2023/cic-iot_2023_3.csv"
5 DATA_PATH_IOT_4 = "/content/drive/MyDrive/Datasets/CIC-IoT 2023/cic-iot_2023_4.csv"
6 DATA_PATH_IOT_5 = "/content/drive/MyDrive/Datasets/CIC-IoT 2023/cic-iot_2023_5.csv"
7
8 df_iot_1 = pd.read_csv(DATA_PATH_IOT_1)
9 df_iot_2 = pd.read_csv(DATA_PATH_IOT_2)
10 df_iot_3 = pd.read_csv(DATA_PATH_IOT_3)
11 df_iot_4 = pd.read_csv(DATA_PATH_IOT_4)
12 df_iot_5 = pd.read_csv(DATA_PATH_IOT_5)
13
14 df_iot_combined = pd.concat([df_iot_1, df_iot_2, df_iot_3, df_iot_4, df_iot_5], ignore_index=True)
15 # Display the first few rows of the dataset
16 df_iot_combined.head()
17

```

	flow_duration	Header_Length	Protocol_Type	Duration	Rate	Srate	Drate	fin_flag_number	syn_flag_number	rst_flag_number
0	0.000000	53.46	5.94	63.36	1.145800	1.145800	0.0	0.0	0.0	1.0
1	0.000000	54.00	6.00	64.00	1.027823	1.027823	0.0	0.0	0.0	0.0
2	2.204616	93.96	6.00	64.00	0.671213	0.671213	0.0	0.0	0.0	1.0
3	0.053618	12497.00	17.00	64.00	47647.897124	47647.897124	0.0	0.0	0.0	0.0
4	0.000000	0.00	1.00	64.00	0.667744	0.667744	0.0	0.0	0.0	0.0

5 rows × 47 columns

```
1 df_iot_combined.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1617141 entries, 0 to 1617140
Data columns (total 47 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   flow_duration    1617141 non-null   float64
 1   Header_Length    1617141 non-null   float64
 2   Protocol_Type    1617141 non-null   float64
 3   Duration          1617141 non-null   float64
 4   Rate              1617141 non-null   float64
 5   Srate             1617141 non-null   float64
 6   Drate             1617141 non-null   float64
 7   fin_flag_number   1617141 non-null   float64
 8   syn_flag_number   1617141 non-null   float64
 9   rst_flag_number   1617141 non-null   float64
 10  psh_flag_number   1617141 non-null   float64
 11  ack_flag_number   1617141 non-null   float64
 12  ece_flag_number   1617141 non-null   float64
 13  cwr_flag_number   1617141 non-null   float64
 14  ack_count         1617141 non-null   float64
 15  syn_count         1617141 non-null   float64
 16  fin_count         1617141 non-null   float64
 17  urg_count         1617141 non-null   float64
 18  rst_count         1617141 non-null   float64
 19  HTTP              1617141 non-null   float64
 20  HTTPS             1617141 non-null   float64
 21  DNS               1617141 non-null   float64
 22  Telnet            1617141 non-null   float64
 23  SMTP              1617141 non-null   float64
 24  SSH               1617141 non-null   float64
 25  IRC               1617141 non-null   float64
 26  TCP               1617141 non-null   float64
 27  UDP               1617141 non-null   float64
 28  DHCP              1617141 non-null   float64
 29  ARP               1617141 non-null   float64
 30  ICMP              1617141 non-null   float64
 31  IPv               1617141 non-null   float64
 32  LLC               1617141 non-null   float64
 33  Tot sum           1617141 non-null   float64
 34  Min               1617141 non-null   float64

```

```

35 Max           1617141 non-null float64
36 AVG           1617141 non-null float64
37 Std           1617141 non-null float64
38 Tot_size      1617141 non-null float64
39 IAT           1617141 non-null float64
40 Number        1617141 non-null float64
41 Magnitue     1617141 non-null float64
42 Radius         1617141 non-null float64
43 Covariance    1617141 non-null float64
44 Variance       1617141 non-null float64
45 Weight         1617141 non-null float64
46 label          1617141 non-null object
dtypes: float64(46), object(1)
memory usage: 579.9+ MB

```

```

1 df = df_iot_combined.copy()
2 print("Shape:", df.shape)
3 print("Columns:", df.columns.tolist()[:8], "...", df.columns.tolist()[-8:])

```

```

Shape: (1617141, 47)
Columns: ['flow_duration', 'Header_Length', 'Protocol Type', 'Duration', 'Rate', 'Srate', 'Drate', 'fin_flag_number'] ... ['IAT', ']

```

```

1 print("Exact duplicate rows:", df.duplicated().sum())
2 nan_cols = df.isna().sum()
3 print("Columns with NaNs:\n", nan_cols[nan_cols>0])
4
5 # Everything is numeric except 'label' (object); confirm:
6 print(df.dtypes.tail(10))
7

```

```

Exact duplicate rows: 0
Columns with NaNs:
Series([], dtype: int64)
Std           float64
Tot_size      float64
IAT           float64
Number        float64
Magnitue      float64
Radius         float64
Covariance    float64
Variance       float64
Weight         float64
label          object
dtype: object

```

✓ Normalize labels into MainClass + Attack (no rows dropped)

```

1 # Clean label strings
2 df["label"] = df["label"].astype(str).str.strip()
3
4 # Subclass → MainClass mapping (only subclasses that actually appear will be used)
5 subclass_mapping_full = {
6     "DDoS": ["DDoS-ICMP_Flood", "DDoS-UDP_Flood", "DDoS-TCP_Flood", "DDoS-PSHACK_Flood",
7               "DDoS-SYN_Flood", "DDoS-RSTFINFlood", "DDoS-SynonymousIP_Flood",
8               "DDoS-ICMP_Fragmentation", "DDoS-UDP_Fragmentation", "DDoS-ACK_Fragmentation",
9               "DDoS-HTTP_Flood", "DDoS-SlowLoris"],
10    "DoS": ["DoS-UDP_Flood", "DoS-TCP_Flood", "DoS-SYN_Flood", "DoS-HTTP_Flood"],
11    "Recon": ["Recon-HostDiscovery", "Recon-OSScan", "Recon-PortScan", "Recon-PingSweep", "VulnerabilityScan"],
12    "Spoofing": ["MITM-ArpSpoofing", "DNS_Spoofing"],
13    "BruteForce": ["DictionaryBruteForce"],
14    "Web-based": ["BrowserHijacking", "XSS", "Uploading_Attack", "SqlInjection", "CommandInjection", "Backdoor_Malware"],
15    "Mirai": ["Mirai-greeth_flood", "Mirai-udplain", "Mirai-greib_flood"],
16    "BENIGN": ["BenignTraffic"]
17 }
18
19 # Restrict to labels that exist in your file (robust to naming diffs)
20 present = set(df["label"].unique())
21 subclass_mapping = {m:[s for s in subs if s in present] for m,subs in subclass_mapping_full.items()}
22 # Drop empty families (no subclass present)
23 subclass_mapping = {m:subs for m,subs in subclass_mapping.items() if len(subs)>0}
24
25 # Build reverse map
26 sub2main = {s:m for m,subs in subclass_mapping.items() for s in subs}
27
28 # New columns

```

```

29 df["MainClass"] = df["label"].map(sub2main).fillna("OTHER_ATTACK")
30 df["Attack"]    = (df["MainClass"]!="BENIGN").astype(int)
31
32 print("Sample:\n", df[["label","MainClass","Attack"]].head())
33 print("\nMainClass counts:\n", df["MainClass"].value_counts())
34 print("\nAttack (0/1):\n", df["Attack"].value_counts())
35

```

Sample:

	label	MainClass	Attack
0	DDoS-SYN_Flood	DDoS	1
1	DDoS-TCP_Flood	DDoS	1
2	DDoS-SynonymousIP_Flood	DDoS	1
3	DDoS-UDP_Flood	DDoS	1
4	DDoS-ICMP_Flood	DDoS	1

MainClass counts:

MainClass	count
DDoS	1177334
DoS	280040
Mirai	91710
BENIGN	37913
Spoofing	16797
Recon	12082
Web-based	787
BruteForce	478

Name: count, dtype: int64

Attack (0/1):

Attack	count
1	1579228
0	37913

Name: count, dtype: int64

✓ Define a proper zero-day protocol

```

1 # Choose unseen subclasses (one per main family) from those present
2 unseen_subclasses = []
3 for main, subs in subclass_mapping.items():
4     if main == "BENIGN":
5         continue
6     vc = df[df["label"].isin(subs)]["label"].value_counts()
7     if vc.empty:
8         continue
9     # Pick least frequent subclass to make generalization harder (deterministic)
10    unseen_subclasses.append(vc.index[-1])
11
12 unseen_subclasses = sorted(set(unseen_subclasses))
13 print("Zero-day (UNSEEN) subclasses:", unseen_subclasses)
14
15 # Partition into SEEN (for tuning) and UNSEEN (for zero-day testing). Benign goes to both.
16 is_unseen = df["label"].isin(unseen_subclasses)
17 df_seen   = df[~is_unseen] | (df["MainClass"]=="BENIGN").copy()
18 df_unseen = df[(is_unseen) | (df["MainClass"]=="BENIGN")].copy()
19
20 print("SEEN shape:", df_seen.shape, "UNSEEN shape:", df_unseen.shape)
21 print("SEEN attack ratio:", df_seen["Attack"].mean(), "UNSEEN attack ratio:", df_unseen["Attack"].mean())
22

```

Zero-day (UNSEEN) subclasses: ['DDoS-SlowLoris', 'DNS_Spoofing', 'DictionaryBruteForce', 'DoS-HTTP_Flood', 'Mirai-greip_flood', 'Re
SEEN shape: (1580977, 49) UNSEEN shape: (74077, 49)
SEEN attack ratio: 0.9760192589772021 UNSEEN attack ratio: 0.4881947163087058

✓ Build splits: Train (benign-only), Validation (seen attacks), Test (unseen attacks)

```

1 from sklearn.model_selection import train_test_split
2
3 # Train = benign-only from SEEN (unsupervised fit benefits from more benign)
4 train_benign = df_seen[df_seen["Attack"]==0].copy()
5
6 # Validation (threshold/hyperparam selection) = mix of benign + SEEN attacks
7 # Keep a sizable, stratified validation subset
8 val_df, _ = train_test_split(

```

```

9     df_seen, test_size=0.70, stratify=df_seen["Attack"], random_state=42
10 )
11
12 # Zero-day test = benign + UNSEEN families (no SEEN attacks)
13 test_df = df_unseen.copy()
14
15 print("Train_benign:", train_benign.shape, "| Val:", val_df.shape, "| Test:", test_df.shape)
16 print("Val attack ratio:", val_df["Attack"].mean(), "Test attack ratio:", test_df["Attack"].mean())
17

```

Train_benign: (37913, 49) | Val: (474293, 49) | Test: (74077, 49)
 Val attack ratio: 0.9760190430809648 Test attack ratio: 0.4881947163087058

✓ Create two test regimes for reporting

```

1 # 1) Realistic test: attacks are rare (e.g., target 10% attack prevalence)
2 target_prev = 0.10
3 ben_test = test_df[test_df["Attack"]==0]
4 atk_test = test_df[test_df["Attack"]==1]
5 n_b = len(ben_test)
6 n_a = min(len(atk_test), int(n_b * target_prev / (1 - target_prev)))
7 atk_down = atk_test.sample(n=n_a, random_state=42) if n_a>0 else atk_test
8 test_realistic = pd.concat([ben_test, atk_down]).sample(frac=1.0, random_state=42).reset_index(drop=True)
9
10 # 2) Stress test: keep original skew from UNSEEN
11 test_stress = test_df.sample(frac=1.0, random_state=42).reset_index(drop=True)
12
13 print("Realistic:", test_realistic.shape, "attack ratio:", test_realistic["Attack"].mean())
14 print("Stress : ", test_stress.shape, "attack ratio:", test_stress["Attack"].mean())
15

```

Realistic: (42125, 49) attack ratio: 0.09998813056379822
 Stress : (74077, 49) attack ratio: 0.4881947163087058

✓ Minimal Exploratory Data Analysis

```

1 def quick_eda(name, dfx):
2     print(f"\n== {name} ==")
3     print("Shape:", dfx.shape, "| Attack ratio:", dfx["Attack"].mean())
4     print("Top MainClass:\n", dfx["MainClass"].value_counts().head(10))
5     # numeric snapshot
6     num_cols = dfx.select_dtypes(include=[np.number]).columns.drop(["Attack"])
7     print("Numeric cols:", len(num_cols))
8     print(dfx[num_cols].describe(percentiles=[.01,.25,.5,.75,.99]).T.head(10))
9
10 quick_eda("Train (benign)", train_benign)
11 quick_eda("Validation (SEEN)", val_df)
12 quick_eda("Test REALISTIC (UNSEEN)", test_realistic)
13 quick_eda("Test STRESS (UNSEEN)", test_stress)
14

```

```

== Train (benign) ==
Shape: (37913, 49) | Attack ratio: 0.0
Top MainClass:
MainClass
BENIGN    37913
Name: count, dtype: int64
Numeric cols: 46
      count        mean         std        min       1%  \
flow_duration  37913.0  3.886023e+01  5.203497e+01  0.000000  0.141755
Header_Length  37913.0  1.020668e+06  1.341055e+06  0.000000  851.092000
Protocol_Type  37913.0  7.470203e+00  2.252745e+00  0.000000  4.800000
Duration       37913.0  1.151078e+02  5.145048e+01  0.000000  50.800000
Rate           37913.0  1.856269e+03  1.731521e+04  0.021591  1.466400
Srate          37913.0  1.856269e+03  1.731521e+04  0.021591  1.466400
Drate          37913.0  0.000000e+00  0.000000e+00  0.000000  0.000000
fin_flag_number 37913.0  0.000000e+00  0.000000e+00  0.000000  0.000000
syn_flag_number 37913.0  5.275235e-05  7.262986e-03  0.000000  0.000000
rst_flag_number 37913.0  0.000000e+00  0.000000e+00  0.000000  0.000000
      25%        50%        75%        99%  \
flow_duration  9.987032  26.472649  5.260494e+01  2.485242e+02
Header_Length  76218.000000  510832.200000  1.402154e+06  6.242186e+06

```

```

Protocol Type      6.000000    6.500000  8.20000e+00  1.48000e+01
Duration          73.600000   99.100000  1.48000e+02  2.32000e+02
Rate              23.852218   52.999576  7.84500e+01  6.995490e+04
Srate             23.852218   52.999576  7.84500e+01  6.995490e+04
Drate             0.000000    0.000000  0.000000e+00  0.000000e+00
fin_flag_number   0.000000    0.000000  0.000000e+00  0.000000e+00
syn_flag_number   0.000000    0.000000  0.000000e+00  0.000000e+00
rst_flag_number   0.000000    0.000000  0.000000e+00  0.000000e+00

      max
flow_duration     1.014878e+03
Header_Length     9.473703e+06
Protocol Type     1.700000e+01
Duration          2.487000e+02
Rate              8.388860e+05
Srate             8.388860e+05
Drate             0.000000e+00
fin_flag_number   0.000000e+00
syn_flag_number   1.000000e+00
rst_flag_number   0.000000e+00

==== Validation (SEEN) ====
Shape: (474293, 49) | Attack ratio: 0.9760190430809648
Top MainClass:
MainClass
DDoS      352838
DoS       83323
Mirai     19596
BENIGN    11374
Recon     3604
Spoofing   3306
Web-based  252
Name: count, dtype: int64
Numeric cols: 46
      count      mean      std      min      1%      25%  \

```

✓ splits to disk

```

1 # Work dirs (adjust drive path if needed)
2 BASE = "/content/drive/MyDrive/colab_zero_day"
3 os.makedirs(f"{BASE}/splits", exist_ok=True)

```

```

1 split_dir = f"{BASE}/splits"
2 train_benign.to_parquet(f"{split_dir}/train_benign_seen.parquet", index=False)
3 val_df.to_parquet(      f"{split_dir}/val_seen.parquet",      index=False)
4 test_realistic.to_parquet(f"{split_dir}/test_unseen_realistic.parquet", index=False)
5 test_stress.to_parquet(   f"{split_dir}/test_unseen_stress.parquet",   index=False)
6
7 meta = {
8     "unseen_subclasses": unseen_subclasses,
9     "target_attack_prevalence_realistic": target_prev,
10    "seed": 42
11 }
12 with open(f"{split_dir}/meta.json", "w") as f:
13     json.dump(meta, f, indent=2)
14
15 print("Saved splits to:", split_dir)
16

```

Saved splits to: /content/drive/MyDrive/colab_zero_day/splits

✓ Preprocessing

```

1 import os, json, math, gc, numpy as np, pandas as pd
2 import matplotlib.pyplot as plt
3 from collections import Counter
4 from scipy import stats
5
6 BASE = "/content/drive/MyDrive/colab_zero_day"
7 EDA_DIR = f"{BASE}/eda"
8 os.makedirs(EDA_DIR, exist_ok=True)
9
10 def savefig(path):
11     plt.tight_layout()

```

```

12     plt.savefig(path, dpi=160, bbox_inches='tight')
13     plt.close()
14
15 def mem_mb(df):
16     return df.memory_usage(deep=True).sum()/1024**2
17
18 # Which columns are meta (not model features)
19 META_COLS = ["Label", "MainClass", "Attack"]
20

```

Basic dataset cards + class balance

```

1 def dataset_card(name, df):
2     print(f"\n== {name} ==")
3     print("shape:", df.shape, "| mem MB:", f"{mem_mb(df):.2f}")
4     print("columns:", len(df.columns))
5     print("attack ratio:", df["Attack"].mean())
6     print("MainClass (top 10):\n", df["MainClass"].value_counts().head(10))
7
8 for nm, d in [("Train (benign-only)", train_benign),
9                 ("Validation (SEEN)", val_df),
10                ("Test REALISTIC (UNSEEN)", test_realistic),
11                ("Test STRESS (UNSEEN)", test_stress)]:
12     dataset_card(nm, d)
13
14 # Bar plots: MainClass counts per split
15 def bar_counts_mainclass(df, title, fname):
16     counts = df["MainClass"].value_counts().sort_values(ascending=False)
17     plt.figure(figsize=(8,4))
18     counts.plot(kind="bar")
19     plt.title(title); plt.ylabel("count"); plt.xlabel("MainClass")
20     savefig(f"{EDA_DIR}/{fname}.png")
21
22 bar_counts_mainclass(val_df, "Validation MainClass distribution", "val_mainclass_counts")
23 bar_counts_mainclass(test_realistic, "Test REALISTIC MainClass distribution", "test_realistic_mainclass_counts")
24 bar_counts_mainclass(test_stress, "Test STRESS MainClass distribution", "test_stress_mainclass_counts")
25
26 # Attack vs Normal counts per split
27 def pie_attack(df, title, fname):
28     plt.figure(figsize=(4,4))
29     vals = df["Attack"].value_counts().reindex([0,1]).fillna(0)
30     plt.pie(vals, labels=["Normal", "Attack"], autopct="%1.1f%%", startangle=90)
31     plt.title(title)
32     savefig(f"{EDA_DIR}/{fname}.png")
33
34 pie_attack(val_df, "Validation Attack vs Normal", "val_attack_pie")
35 pie_attack(test_realistic, "Test REALISTIC Attack vs Normal", "test_realistic_attack_pie")
36 pie_attack(test_stress, "Test STRESS Attack vs Normal", "test_stress_attack_pie")
37

```

```

==== Train (benign-only) ====
shape: (37913, 49) | mem MB: 18.11
columns: 49
attack ratio: 0.0
MainClass (top 10):
    MainClass
    BENIGN      37913
Name: count, dtype: int64

==== Validation (SEEN) ====
shape: (474293, 49) | mem MB: 226.78
columns: 49
attack ratio: 0.9760190430809648
MainClass (top 10):
    MainClass
    DDoS        352838
    DoS         83323
    Mirai       19596
    BENIGN      11374
    Recon        3604
    Spoofing     3306
    Web-based    252
Name: count, dtype: int64

==== Test REALISTIC (UNSEEN) ====

```

```

shape: (42125, 49) | mem MB: 19.82
columns: 49
attack ratio: 0.09998813056379822
MainClass (top 10):
    MainClass
BENIGN      37913
Mirai        3035
Spoofing     718
DoS          271
DDoS         103
BruteForce   67
Recon         13
Web-based    5
Name: count, dtype: int64

==== Test STRESS (UNSEEN) ====
shape: (74077, 49) | mem MB: 34.91
columns: 49
attack ratio: 0.4881947163087058
MainClass (top 10):
    MainClass
BENIGN      37913
Mirai        26158
Spoofing     6141
DoS          2495
DDoS         791
BruteForce   478
Recon         69
Web-based    32
Name: count, dtype: int64

```

▼ Numeric summary & missingness

```

1 def numeric_columns(df):
2     return df.select_dtypes(include=[np.number]).columns.difference(META_COLS)
3
4 num_cols_all = numeric_columns(val_df) # use val_df to define feature set
5 print("Numeric feature count:", len(num_cols_all))
6
7 # Missingness table (should be tiny; if not, we'll impute later)
8 miss_tbl = pd.DataFrame({
9     "train_benign": train_benign[num_cols_all].isna().sum(),
10    "val": val_df[num_cols_all].isna().sum(),
11    "test_realistic": test_realistic[num_cols_all].isna().sum(),
12    "test_stress": test_stress[num_cols_all].isna().sum(),
13 })
14 miss_tbl.to_csv(f"{EDA_DIR}/missingness.csv")
15 print("Missingness saved -> missingness.csv")
16
17 # Quick numeric stats (validation)
18 desc = val_df[num_cols_all].describe(percentiles=[.01,.05,.25,.5,.75,.95,.99]).T
19 desc.to_csv(f"{EDA_DIR}/val_numeric_describe.csv")
20 desc.head()
21

```

Numeric feature count: 46
Missingness saved -> missingness.csv

	count	mean	std	min	1%	5%	25%	50%	75%	95%	99%	max
ARP	474293.0	0.000038	0.006160	0.0	0.0	0.0	0.0	0.0	0.000000	0.000000	0.000000	1.000000e+00
AVG	474293.0	115.785542	234.994619	42.0	42.0	42.0	50.0	54.0	54.015000	587.441335	1004.797783	1.165000e+04
Covariance	474293.0	30103.214257	358113.590216	0.0	0.0	0.0	0.0	0.0	0.910903	25292.212302	587493.689942	1.076490e+08
DHCP	474293.0	0.000002	0.001452	0.0	0.0	0.0	0.0	0.0	0.000000	0.000000	0.000000	1.000000e+00
DNS	474293.0	0.000093	0.009631	0.0	0.0	0.0	0.0	0.0	0.000000	0.000000	0.000000	1.000000e+00

▼ Feature separation ranking (KS & Cohen's d)

```

1 def feature_separation(df, cols, target="Attack"):
2     rows=[]
3     a = df[df[target]==1][cols]
4     n = df[df[target]==0][cols]

```

```

5   for c in cols:
6       # Two-sample KS statistic (distributional difference)
7       ks_stat, ks_p = stats.ks_2samp(a[c].values, n[c].values, alternative='two-sided', mode='auto')
8       # Cohen's d (mean difference / pooled sd)
9       mu1, mu0 = a[c].mean(), n[c].mean()
10      s1, s0 = a[c].std(ddof=1), n[c].std(ddof=1)
11      n1, n0 = a[c].shape[0], n[c].shape[0]
12      sp = math.sqrt((n1-1)*s1**2 + (n0-1)*s0**2) / max(n1+n0-2,1)
13      d = (mu1 - mu0) / (sp + 1e-12)
14      rows.append((c, ks_stat, ks_p, d, mu0, mu1))
15  out = pd.DataFrame(rows, columns=["feature", "KS", "KS_p", "Cohen_d", "mean_normal", "mean_attack"])
16  out["abs_d"] = out["Cohen_d"].abs()
17  out.sort_values(["KS", "abs_d"], ascending=[False, False], inplace=True)
18  return out
19
20 sep_tbl = feature_separation(val_df, num_cols_all)
21 sep_tbl.to_csv(f"{EDA_DIR}/feature_separation_val.csv", index=False)
22 sep_tbl.head(10)
23

```

	feature	KS	KS_p	Cohen_d	mean_normal	mean_attack	abs_d
41	rst_count	0.970092	0.0	-3.734265	1079.206057	12.522938	3.734265
45	urg_count	0.959992	0.0	-1.654283	117.357471	3.321979	1.654283
31	Variance	0.894642	0.0	-4.134234	0.864167	0.072552	4.134234
15	Magnitue	0.884271	0.0	-2.324234	30.265314	12.288978	2.324234
28	Tot size	0.883263	0.0	-2.422983	636.306384	103.054970	2.422983
1	AVG	0.882726	0.0	-2.410241	634.458300	103.041663	2.410241
25	Std	0.861164	0.0	-3.402065	500.857297	20.361353	3.402065
20	Radius	0.860759	0.0	-3.397181	707.597746	28.781234	3.397181
16	Max	0.851222	0.0	-3.521473	1743.614360	131.890664	3.521473
2	Covariance	0.838170	0.0	-1.890230	664782.049875	14509.045350	1.890230

▼ Histograms for the top 12 discriminative features

```

1 TOPK = 12
2 top_feats = sep_tbl["feature"].head(TOPK).tolist()
3 print("Top features by separation:", top_feats)
4
5 def plot_hist_by_class(df, cols, title_prefix, fname_prefix):
6     for c in cols:
7         plt.figure(figsize=(5,3))
8         # Normal
9         x0 = df[df["Attack"]==0][c].values
10        x1 = df[df["Attack"]==1][c].values
11        # Clip extremes for viz (1st-99th percentile)
12        lo = np.nanpercentile(df[c].values, 1)
13        hi = np.nanpercentile(df[c].values, 99)
14        bins = 50
15        plt.hist(np.clip(x0, lo, hi), bins=bins, alpha=0.6, label="Normal", density=True)
16        plt.hist(np.clip(x1, lo, hi), bins=bins, alpha=0.6, label="Attack", density=True)
17        plt.xlabel(c); plt.ylabel("density")
18        plt.title(f"{title_prefix}: {c}")
19        plt.legend()
20        savefig(f'{EDA_DIR}/{fname_prefix}_{c.replace(' ', '_)}.png')
21
22 plot_hist_by_class(val_df, top_feats, "Validation distributions", "val_hist")
23

```

Top features by separation: ['rst_count', 'urg_count', 'Variance', 'Magnitue', 'Tot size', 'AVG', 'Std', 'Radius', 'Max', 'Covarian

▼ Correlation heatmap + highly correlated pairs

```

1 # Correlation on validation (only numeric)
2 corr = val_df[num_cols_all].corr().fillna(0.0)

```

```

3 plt.figure(figsize=(8,6))
4 plt.imshow(corr.values, aspect='auto', interpolation='nearest')
5 plt.colorbar(label="Pearson r")
6 plt.xticks(range(len(num_cols_all)), num_cols_all, rotation=90, fontsize=6)
7 plt.yticks(range(len(num_cols_all)), num_cols_all, fontsize=6)
8 plt.title("Validation correlation heatmap")
9 savefig(f"{EDA_DIR}/corr_heatmap_val.png")
10
11 # Extract high-corr pairs
12 pairs = []
13 thr = 0.98
14 for i in range(len(num_cols_all)):
15     for j in range(i+1, len(num_cols_all)):
16         r = corr.iat[i,j]
17         if abs(r) >= thr:
18             pairs.append((num_cols_all[i], num_cols_all[j], r))
19 high_corr_pairs = sorted(pairs, key=lambda x: -abs(x[2]))
20 pd.DataFrame(high_corr_pairs, columns=["feat_a","feat_b","r"]).to_csv(f"{EDA_DIR}/high_corr_pairs.csv", index=False)
21 print("High corr pairs (|r|>=0.98):", len(high_corr_pairs))
22

```

High corr pairs (|r|>=0.98): 7

Outlier analysis & tail heaviness

```

1 def tail_heaviness(df, cols):
2     rows=[]
3     for c in cols:
4         x = df[c].values
5         # robust metrics
6         q1,q3 = np.nanpercentile(x, [25,75]); iqr = q3 - q1
7         p01,p99 = np.nanpercentile(x, [1,99])
8         kurt = stats.kurtosis(x, fisher=True, nan_policy='omit')
9         skw = stats.skew(x, nan_policy='omit')
10        out_frac = np.mean((x < q1 - 3*iqr) | (x > q3 + 3*iqr))
11        rows.append((c, skw, kurt, out_frac, p01, p99, q1, q3))
12    T = pd.DataFrame(rows, columns=["feature","skew","kurtosis","outlier_frac","p01","p99","q1","q3"])
13    T.sort_by = "outlier_frac"
14    return T.sort_values("outlier_frac", ascending=False)
15
16 tails = tail_heaviness(val_df, num_cols_all)
17 tails.to_csv(f"{EDA_DIR}/tail_heaviness_val.csv", index=False)
18 tails.head(10)
19
20 # Plot outlier rate bar for top 20
21 plt.figure(figsize=(8,4))
22 top_out = tails.head(20)
23 plt.bar(range(len(top_out)), top_out["outlier_frac"])
24 plt.xticks(range(len(top_out)), top_out["feature"], rotation=90, fontsize=7)
25 plt.ylabel("fraction beyond 3*IQR")
26 plt.title("Top-20 heavy-tail/outlier features (validation)")
27 savefig(f"{EDA_DIR}/outlier_frac_bar.png")
28
29 # Heuristic recommendation
30 heavy_tail_frac = (tails["outlier_frac"] > 0.01).mean() # >1% extreme
31 suggest_scaler = "RobustScaler" if heavy_tail_frac > 0.1 else "StandardScaler"
32 print(f"Suggested scaler based on tails: {suggest_scaler} (heavy-tail features: {heavy_tail_frac*100:.1f}% )")
33

```

Suggested scaler based on tails: RobustScaler (heavy-tail features: 71.7%)

Constant / near-constant features, zero variance

```

1 def constant_features(df, cols, thresh_unique=1):
2     nun = df[cols].nunique(dropna=False)
3     const = nun[nun <= thresh_unique].index.tolist()
4     return const
5
6 const_feats = constant_features(train_benign, num_cols_all, 1)
7 print("Constant features in TRAIN (drop):", const_feats)
8 pd.Series(const_feats).to_csv(f"{EDA_DIR}/constant_features.csv", index=False)
9

```

```
Constant features in TRAIN (drop): ['DHCP', 'Drate', 'IRC', 'SMTP', 'SSH', 'Telnet', 'cwr_flag_number', 'ece_flag_number', 'fin fla
```

✓ Save a preprocessing plan for the next step

```

1 # Pick which feature to keep from each high-corr pair (keep the one with higher separation)
2 sep = sep_tbl.set_index("feature")
3 drop_due_corr = []
4 keep_due_corr = set()
5
6 for a,b,r in high_corr_pairs:
7     if a in drop_due_corr or b in drop_due_corr:
8         continue
9     # Compare absolute Cohen's d (fallback to KS)
10    da = abs(sep.loc[a, "Cohen_d"]) if a in sep.index else 0.0
11    db = abs(sep.loc[b, "Cohen_d"]) if b in sep.index else 0.0
12    if da >= db:
13        drop_due_corr.append(b); keep_due_corr.add(a)
14    else:
15        drop_due_corr.append(a); keep_due_corr.add(b)
16
17 drop_due_corr = sorted(set(drop_due_corr))
18 print("Drop due to high correlation:", drop_due_corr[:15], f"... total={len(drop_due_corr)}")
19 pd.Series(drop_due_corr).to_csv(f"{EDA_DIR}/drop_due_to_corr.csv", index=False)
20
21 # Tail clipping suggestions (top 10 heaviest)
22 clip_suggestions = tails.head(10)[["feature", "p01", "p99"]].to_dict(orient="records")
23 with open(f"{EDA_DIR}/clip_suggestions.json", "w") as f:
24     json.dump(clip_suggestions, f, indent=2)
25
26 preproc_plan = {
27     "suggested_scaler": suggest_scaler,
28     "constant_features_drop": const_feats,
29     "high_corr_drop": drop_due_corr,
30     "top_sep_features": top_feats,
31     "clip_suggestions": clip_suggestions,
32     "feature_list_candidate": [c for c in num_cols_all if c not in set(const_feats) | set(drop_due_corr)]
33 }
34 with open(f"{EDA_DIR}/preprocessing_plan.json", "w") as f:
35     json.dump(preproc_plan, f, indent=2)
36
37 print("Saved preprocessing plan -> preprocessing_plan.json")
38

```

```
Drop due to high correlation: ['IAT', 'LLC', 'Radius', 'Srate', 'Weight', 'ack_count'] ... total=6
Saved preprocessing plan -> preprocessing_plan.json
```

✓ Build the preprocessing transformer (winsorize + scale)

```

1 import os, json, numpy as np, pandas as pd
2 from sklearn.base import BaseEstimator, TransformerMixin
3 from sklearn.preprocessing import StandardScaler, RobustScaler
4 from sklearn.pipeline import Pipeline
5
6 BASE = "/content/drive/MyDrive/colab_zero_day"
7 PLAN_PATH = f"{BASE}/eda/preprocessing_plan.json"
8
9 with open(PLAN_PATH, "r") as f:
10     preproc_plan = json.load(f)
11
12 features_keep = preproc_plan["feature_list_candidate"] # numeric features we'll keep
13 clip_suggestions = preproc_plan.get("clip_suggestions", [])
14 scaler_choice = preproc_plan.get("suggested_scaler", "RobustScaler")
15
16 # Map clip targets -> we will recompute percentiles from TRAIN BENIGN to avoid leakage
17 clip_targets = [d["feature"] for d in clip_suggestions]
18
19 class ColumnClipper(BaseEstimator, TransformerMixin):
20     """
21         Winsorize specified columns using percentiles computed on fit(X).
22         thresholds_: dict {col: (lo, hi)}

```

```

23 """
24 def __init__(self, cols, lower=1.0, upper=99.0):
25     self.cols = list(cols)
26     self.lower = lower
27     self.upper = upper
28     self.thresholds_ = {}
29
30 def fit(self, X, y=None):
31     df = pd.DataFrame(X, columns=self.feature_names_in_) if hasattr(self, "feature_names_in_") else X
32     # If X is DataFrame, use its columns; else assume previously set
33     if isinstance(df, pd.DataFrame):
34         for c in self.cols:
35             if c in df.columns:
36                 lo = np.nanpercentile(df[c].values, self.lower)
37                 hi = np.nanpercentile(df[c].values, self.upper)
38                 self.thresholds_[c] = (lo, hi)
39     return self
40
41 def transform(self, X):
42     if isinstance(X, pd.DataFrame):
43         df = X.copy()
44         for c,(lo,hi) in self.thresholds_.items():
45             if c in df.columns:
46                 x = df[c].values
47                 x = np.clip(x, lo, hi)
48                 df[c] = x
49     return df
50 else:
51     # If numpy array was passed, we need columns; prefer passing DataFrames
52     return X
53
54 class ColumnSelector(BaseEstimator, TransformerMixin):
55     def __init__(self, cols):
56         self.cols = list(cols)
57     def fit(self, X, y=None):
58         return self
59     def transform(self, X):
60         return X[self.cols].copy() # keep DataFrame for named columns
61
62 # Choose scaler
63 Scaler = RobustScaler if scaler_choice == "RobustScaler" else StandardScaler
64
65 # Build preprocessing pipeline:
66 # 1) select columns
67 # 2) winsorize tails for selected heavy-tail features (recomputed from train_benign)
68 # 3) scale
69 preprocess = Pipeline(steps=[
70     ("select", ColumnSelector(features_keep)),
71     ("winsor", ColumnClipper(cols=clip_targets, lower=1.0, upper=99.0)),
72     ("scale", Scaler(with_centering=True, with_scaling=True))
73 ])
74
75 # Fit on BENIGN TRAIN ONLY (no leakage)
76 preprocess.fit(train_benign)
77
78 # Transform all splits
79 X_train = preprocess.transform(train_benign)      # benign-only
80 X_val   = preprocess.transform(val_df)
81 X_tr    = preprocess.transform(test_realistic)
82 X_ts    = preprocess.transform(test_stress)
83
84 y_val = val_df["Attack"].astype(int).values
85 y_tr  = test_realistic["Attack"].astype(int).values
86 y_ts  = test_stress["Attack"].astype(int).values
87
88 print("Shapes -> X_train:", X_train.shape, "X_val:", X_val.shape, "X_tr:", X_tr.shape, "X_ts:", X_ts.shape)
89

```

Shapes -> X_train: (37913, 30) X_val: (474293, 30) X_tr: (42125, 30) X_ts: (74077, 30)

- ✓ Thresholding utilities & metrics (F2 and FPR-controlled)

```
1 import numpy as np
2 from sklearn.metrics import precision_recall_curve, classification_report, roc_auc_score, confusion_matrix, roc_curve, precision
```

```

3
4 def pick_thresh_by_Fbeta(y_true, scores, beta=2.0):
5     P, R, thr = precision_recall_curve(y_true, scores)
6     P, R = P[:-1], R[:-1] # last point has no threshold
7     fbeta = (1+beta**2) * (P*R) / (beta**2 * P + R + 1e-12)
8     i = np.nanargmax(fbeta)
9     return float(thr[i]), float(P[i]), float(R[i]), float(fbeta[i])
10
11 def pick_thresh_at_fpr(y_true, scores, max_fpr=0.05):
12     fpr, tpr, thr = roc_curve(y_true, scores)
13     ok = np.where(fpr <= max_fpr)[0]
14     if len(ok)==0:
15         j = np.argmax(tpr - fpr)
16         return float(thr[j]), float(fpr[j]), float(tpr[j])
17     i = ok[np.argmax(tpr[ok])]
18     return float(thr[i]), float(fpr[i]), float(tpr[i])
19
20 def summarize(y_true, y_pred, scores, name=""):
21     P,R,F1,_ = precision_recall_fscore_support(y_true, y_pred, average='binary', zero_division=0)
22     auc_ = roc_auc_score(y_true, scores) if len(np.unique(y_true))>1 else np.nan
23     cm = confusion_matrix(y_true, y_pred)
24     print(f"\n{name}:\n")
25     print(f"Precision={P:.4f} Recall={R:.4f} F1={F1:.4f} ROC-AUC={auc_:.4f}")
26     print("Confusion matrix:\n", cm)
27     print(classification_report(y_true, y_pred, digits=4))
28

```

▼ Train & evaluate Isolation Forest

```

1 from sklearn.ensemble import IsolationForest
2
3 iforest = IsolationForest(
4     n_estimators=400, max_samples='auto',
5     contamination='auto', # we will set threshold from validation
6     random_state=42, n_jobs=-1
7 ).fit(X_train) # benign only
8
9 # Anomaly scores (higher = more anomalous)
10 val_scores = -iforest.score_samples(X_val)
11 tr_scores = -iforest.score_samples(X_tr)
12 ts_scores = -iforest.score_samples(X_ts)
13
14 # Strategy A: F2-optimized threshold (recall-oriented)
15 th_f2, p_f2, r_f2, f2 = pick_thresh_by_Fbeta(y_val, val_scores, beta=2.0)
16 y_tr_f2 = (tr_scores >= th_f2).astype(int)
17 y_ts_f2 = (ts_scores >= th_f2).astype(int)
18
19 print(f"IForest threshold F2-opt: {th_f2:.6f} (val P={p_f2:.3f}, R={r_f2:.3f}, F2={f2:.3f})")
20 summarize(y_tr, y_tr_f2, tr_scores, "IForest | REALISTIC | F2-opt")
21 summarize(y_ts, y_ts_f2, ts_scores, "IForest | STRESS | F2-opt")
22
23 # Strategy B: control FPR on validation (e.g., <= 5%)
24 th_fpr, fpr_sel, tpr_sel = pick_thresh_at_fpr(y_val, val_scores, max_fpr=0.05)
25 y_tr_fpr = (tr_scores >= th_fpr).astype(int)
26 y_ts_fpr = (ts_scores >= th_fpr).astype(int)
27
28 print(f"IForest threshold FPR<=5%: {th_fpr:.6f} (val FPR={fpr_sel:.3f}, TPR={tpr_sel:.3f})")
29 summarize(y_tr, y_tr_fpr, tr_scores, "IForest | REALISTIC | FPR<=5%")
30 summarize(y_ts, y_ts_fpr, ts_scores, "IForest | STRESS | FPR<=5%")
31

```

IForest threshold F2-opt: 0.384001 (val P=0.978, R=1.000, F2=0.995)

== IForest | REALISTIC | F2-opt ==

Precision=0.1059 Recall=0.9945 F1=0.1914 ROC-AUC=0.9103

Confusion matrix:

[[2538 35375]	[23 4189]]
----------------	-------------

	precision	recall	f1-score	support
0	0.9910	0.0669	0.1254	37913
1	0.1059	0.9945	0.1914	4212
accuracy			0.1597	42125
macro avg	0.5484	0.5307	0.1584	42125

```

weighted avg    0.9025    0.1597    0.1320    42125

== IForest | STRESS      | F2-opt ==
Precision=0.5045  Recall=0.9960  F1=0.6698  ROC-AUC=0.9129
Confusion matrix:
[[ 2538 35375]
 [ 144 36020]]
precision      recall   f1-score   support
0            0.9463    0.0669    0.1250    37913
1            0.5045    0.9960    0.6698    36164

accuracy          0.5205
macro avg       0.7254    0.5315    0.3974    74077
weighted avg     0.7306    0.5205    0.3910    74077

IForest threshold FPR<=5%: 0.525049 (val FPR=0.050, TPR=0.378)

== IForest | REALISTIC | FPR<=5% ==
Precision=0.6353  Recall=0.7602  F1=0.6922  ROC-AUC=0.9103
Confusion matrix:
[[36075 1838]
 [ 1010 3202]]
precision      recall   f1-score   support
0            0.9728    0.9515    0.9620    37913
1            0.6353    0.7602    0.6922    4212

accuracy          0.9324
macro avg       0.8040    0.8559    0.8271    42125
weighted avg     0.9390    0.9324    0.9350    42125

== IForest | STRESS      | FPR<=5% ==
Precision=0.9374  Recall=0.7612  F1=0.8402  ROC-AUC=0.9129
Confusion matrix:
[[36075 1838]
 [ 8636 27528]]
precision      recall   f1-score   support
0            0.8068    0.9515    0.8732    37913
1            0.9374    0.7612    0.8402    36164

```

▼ Train & evaluate Autoencoder

```

1 import tensorflow as tf
2 from tensorflow import keras
3 tf.random.set_seed(42)
4
5 input_dim = X_train.shape[1]
6 inp = keras.Input(shape=(input_dim,))
7 x = keras.layers.Dense(256, activation='relu')(inp)
8 x = keras.layers.Dense(128, activation='relu')(x)
9 z = keras.layers.Dense(64,  activation='relu')(x)
10 x = keras.layers.Dense(128, activation='relu')(z)
11 x = keras.layers.Dense(256, activation='relu')(x)
12 out = keras.layers.Dense(input_dim, activation='linear')(x)
13
14 ae = keras.Model(inp, out)
15 ae.compile(optimizer=keras.optimizers.Adam(1e-3), loss='mse')
16 es = keras.callbacks.EarlyStopping(monitor='loss', patience=5, restore_best_weights=True)
17 hist = ae.fit(X_train, X_train, epochs=60, batch_size=1024, shuffle=True, callbacks=[es], verbose=0)
18
19 # Reconstruction error = anomaly score
20 val_rec = ae.predict(X_val, batch_size=4096, verbose=0)
21 tr_rec = ae.predict(X_tr, batch_size=4096, verbose=0)
22 ts_rec = ae.predict(X_ts, batch_size=4096, verbose=0)
23
24 val_err = np.mean((X_val - val_rec)**2, axis=1)
25 tr_err = np.mean((X_tr - tr_rec )**2, axis=1)
26 ts_err = np.mean((X_ts - ts_rec )**2, axis=1)
27
28 # Thresholds
29 th_f2_ae, p_f2_ae, r_f2_ae, f2_ae = pick_thresh_by_Fbeta(y_val, val_err, beta=2.0)
30 y_tr_f2_ae = (tr_err >= th_f2_ae).astype(int)
31 y_ts_f2_ae = (ts_err >= th_f2_ae).astype(int)
32

```

```

33 print(f"AE threshold F2-opt: {th_f2_ae:.6f} (val P={p_f2_ae:.3f}, R={r_f2_ae:.3f}, F2={f2_ae:.3f})")
34 summarize(y_tr, y_tr_f2_ae, tr_err, "AE | REALISTIC | F2-opt")
35 summarize(y_ts, y_ts_f2_ae, ts_err, "AE | STRESS | F2-opt")
36
37 th_fpr_ae, fpr_sel_ae, tpr_sel_ae = pick_thresh_at_fpr(y_val, val_err, max_fpr=0.05)
38 y_tr_fpr_ae = (tr_err >= th_fpr_ae).astype(int)
39 y_ts_fpr_ae = (ts_err >= th_fpr_ae).astype(int)
40
41 print(f"AE threshold FPR<=5%: {th_fpr_ae:.6f} (val FPR={fpr_sel_ae:.3f}, TPR={tpr_sel_ae:.3f})")
42 summarize(y_tr, y_tr_fpr_ae, tr_err, "AE | REALISTIC | FPR<=5%")
43 summarize(y_ts, y_ts_fpr_ae, ts_err, "AE | STRESS | FPR<=5%")
44

```

AE threshold F2-opt: 0.002784 (val P=0.982, R=0.999, F2=0.995)

== AE | REALISTIC | F2-opt ==
Precision=0.1258 Recall=0.9919 F1=0.2232 ROC-AUC=0.9312
Confusion matrix:
[[8868 29045]
[34 4178]]

	precision	recall	f1-score	support
0	0.9962	0.2339	0.3789	37913
1	0.1258	0.9919	0.2232	4212

accuracy 0.3097 42125
macro avg 0.5610 0.6129 0.3010 42125
weighted avg 0.9091 0.3097 0.3633 42125

== AE | STRESS | F2-opt ==
Precision=0.5521 Recall=0.9899 F1=0.7088 ROC-AUC=0.9317
Confusion matrix:
[[8868 29045]
[367 35797]]

	precision	recall	f1-score	support
0	0.9603	0.2339	0.3762	37913
1	0.5521	0.9899	0.7088	36164

accuracy 0.6030 74077
macro avg 0.7562 0.6119 0.5425 74077
weighted avg 0.7610 0.6030 0.5386 74077

AE threshold FPR<=5%: 0.054767 (val FPR=0.050, TPR=0.562)

== AE | REALISTIC | FPR<=5% ==
Precision=0.6629 Recall=0.8170 F1=0.7319 ROC-AUC=0.9312
Confusion matrix:
[[36163 1750]
[771 3441]]

	precision	recall	f1-score	support
0	0.9791	0.9538	0.9663	37913
1	0.6629	0.8170	0.7319	4212

accuracy 0.9402 42125
macro avg 0.8210 0.8854 0.8491 42125
weighted avg 0.9475 0.9402 0.9429 42125

== AE | STRESS | FPR<=5% ==
Precision=0.9441 Recall=0.8172 F1=0.8761 ROC-AUC=0.9317
Confusion matrix:
[[36163 1750]
[6610 29554]]

	precision	recall	f1-score	support
0	0.8455	0.9538	0.8964	37913
1	0.9441	0.8172	0.8761	36164

✓ Save artifacts (preprocessor, models, thresholds)

```

1 import joblib, json, os
2 ART = f"{BASE}/models"; os.makedirs(ART, exist_ok=True)
3
4 joblib.dump(preprocess, f"{ART}/preprocess.joblib")
5 joblib.dump(iforest, f"{ART}/iforest.joblib")

```

```

6 ae.save(f"{ART}/autoencoder_ae.keras")
7
8 thresholds = {
9     "iforest": {"F2": float(th_f2), "FPR5": float(th_fpr)},
10    "autoencoder": {"F2": float(th_f2_ae), "FPR5": float(th_fpr_ae)}
11 }
12 with open(f"{ART}/thresholds.json", "w") as f:
13     json.dump(thresholds, f, indent=2)
14
15 print("Saved ->", ART)
16

```

Saved -> /content/drive/MyDrive/colab_zero_day/models

```

1 import numpy as np, pandas as pd
2 from sklearn.metrics import average_precision_score, precision_recall_curve
3
4 def pr_summary(y_true, scores, name=""):
5     ap = average_precision_score(y_true, scores)
6     P,R,Thr = precision_recall_curve(y_true, scores)
7     print(f"{name} PR-AUC(AP) = {ap:.4f}")
8     return ap
9
10 # Use your already computed scores: IForest (ts_scores) and AE (ts_err)
11 ap_if = pr_summary(y_ts, ts_scores, "IForest | STRESS")
12 ap_ae = pr_summary(y_ts, ts_err, "AE | STRESS")
13
14 # Per-MainClass recall on UNSEEN test at your FPR<=5% thresholds
15 def per_family_recall(df_test, y_pred, family_col="MainClass"):
16     T = df_test[[family_col,"Attack"]].copy()
17     T["pred"] = y_pred
18     # Only attacks; benign don't have 'family'
19     atk = T[T["Attack"]==1]
20     rec = atk.groupby(family_col).apply(lambda g: (g["pred"]==1).mean()).sort_values(ascending=False)
21     return rec
22
23 rec_if = per_family_recall(test_stress, (ts_scores >= th_fpr).astype(int))
24 rec_ae = per_family_recall(test_stress, (ts_err >= th_fpr_ae).astype(int))
25 print("\nPer-family recall (IForest, FPR<=5%):\n", rec_if)
26 print("\nPer-family recall (AE, FPR<=5%):\n", rec_ae)
27

```

IForest | STRESS PR-AUC(AP) = 0.8567
AE | STRESS PR-AUC(AP) = 0.9048

Per-family recall (IForest, FPR<=5%):

MainClass	Value
Mirai	0.949881
DoS	0.473347
Spoofing	0.234815
DDoS	0.059418
Recon	0.028986
BruteForce	0.018828
Web-based	0.000000

dtype: float64

Per-family recall (AE, FPR<=5%):

MainClass	Value
Mirai	0.981382
DoS	0.783166
DDoS	0.335019
Recon	0.260870
Spoofing	0.260707
BruteForce	0.089958
Web-based	0.062500

dtype: float64

```

1 # Rank-normalize scores to [0,1] then average (robust across scales)
2 def rank_norm(x):
3     r = pd.Series(x).rank(method="average").values
4     return (r - r.min()) / (r.max() - r.min() + 1e-12)
5
6 val_ens = 0.5*rank_norm(val_scores) + 0.5*rank_norm(val_err)
7 tr_ens = 0.5*rank_norm(tr_scores) + 0.5*rank_norm(tr_err)
8 ts_ens = 0.5*rank_norm(ts_scores) + 0.5*rank_norm(ts_err)
9
10 # Pick threshold on validation (again FPR<=5% and F2)
11 th_ens_f2, _, _, _ = pick_thresh_by_Fbeta(y_val, val_ens, beta=2.0)

```

```

12 th_ens_fpr, _, _ = pick_thresh_at_fpr(y_val, val_ens, max_fpr=0.05)
13
14 from sklearn.metrics import roc_auc_score
15 y_ts_f2 = (ts_ens >= th_ens_f2).astype(int)
16 y_ts_fpr5 = (ts_ens >= th_ens_fpr).astype(int)
17
18 summarize(y_ts, y_ts_f2, ts_ens, "Ensemble | STRESS | F2-opt")
19 summarize(y_ts, y_ts_fpr5, ts_ens, "Ensemble | STRESS | FPR<=5%")
20
21 # Also report REALISTIC
22 y_tr_f2 = (tr_ens >= th_ens_f2).astype(int)
23 y_tr_fpr5 = (tr_ens >= th_ens_fpr).astype(int)
24 summarize(y_tr, y_tr_f2, tr_ens, "Ensemble | REALISTIC | F2-opt")
25 summarize(y_tr, y_tr_fpr5, tr_ens, "Ensemble | REALISTIC | FPR<=5%")
26

```

== Ensemble | STRESS | F2-opt ==
Precision=0.4891 Recall=0.9999 F1=0.6569 ROC-AUC=0.9308

Confusion matrix:

```
[[ 135 37778]
 [ 2 36162]]
```

	precision	recall	f1-score	support
0	0.9854	0.0036	0.0071	37913
1	0.4891	0.9999	0.6569	36164

	accuracy			
macro avg	0.7372	0.5018	0.3320	74077
weighted avg	0.7431	0.4900	0.3243	74077

== Ensemble | STRESS | FPR<=5% ==
Precision=0.8666 Recall=0.8893 F1=0.8778 ROC-AUC=0.9308

Confusion matrix:

```
[[32963 4950]
 [ 4003 32161]]
```

	precision	recall	f1-score	support
0	0.8917	0.8694	0.8804	37913
1	0.8666	0.8893	0.8778	36164

	accuracy			
macro avg	0.8792	0.8794	0.8791	74077
weighted avg	0.8795	0.8791	0.8792	74077

== Ensemble | REALISTIC | F2-opt ==
Precision=0.1001 Recall=1.0000 F1=0.1821 ROC-AUC=0.9331

Confusion matrix:

```
[[ 67 37846]
 [ 0 4212]]
```

	precision	recall	f1-score	support
0	1.0000	0.0018	0.0035	37913
1	0.1001	1.0000	0.1821	4212

	accuracy			
macro avg	0.5501	0.5009	0.0928	42125
weighted avg	0.9100	0.1016	0.0214	42125

== Ensemble | REALISTIC | FPR<=5% ==
Precision=0.1970 Recall=0.9577 F1=0.3269 ROC-AUC=0.9331

Confusion matrix:

```
[[21475 16438]
 [ 178 4034]]
```

	precision	recall	f1-score	support
0	0.9918	0.5664	0.7210	37913
1	0.1970	0.9577	0.3269	4212

	accuracy			
macro avg	0.5944	0.7621	0.5240	42125

>Loading UNSW-2015 Dataset

```

1 import pandas as pd
2 NB15_1 = pd.read_csv('/content/drive/MyDrive/Datasets/UNSW-2015/UNSW-NB15_1.csv')

```

```
3 NB15_2 = pd.read_csv('/content/drive/MyDrive/Datasets/UNSW-2015/UNSW-NB15_2.csv')
4 NB15_3 = pd.read_csv('/content/drive/MyDrive/Datasets/UNSW-2015/UNSW-NB15_3.csv')
5 NB15_4 = pd.read_csv('/content/drive/MyDrive/Datasets/UNSW-2015/UNSW-NB15_4.csv')
6 NB15_features = pd.read_csv('/content/drive/MyDrive/Datasets/UNSW-2015/NUSW-NB15_features.csv', encoding='cp1252')

/tmp/ipython-input-1562369842.py:2: DtypeWarning: Columns (1,3,47) have mixed types. Specify dtype option on import or set low_memo
NB15_1 = pd.read_csv('/content/drive/MyDrive/Datasets/UNSW-2015/UNSW-NB15_1.csv')
/tmp/ipython-input-1562369842.py:3: DtypeWarning: Columns (3,39,47) have mixed types. Specify dtype option on import or set low_memo
NB15_2 = pd.read_csv('/content/drive/MyDrive/Datasets/UNSW-2015/UNSW-NB15_2.csv')
```

```
1 NB15_features
```


No.	Name	Type	Description
0	1	srcip	nominal Source IP address
1	2	sport	integer Source port number
2	3	dstip	nominal Destination IP address
3	4	dsport	integer Destination port number
4	5	proto	nominal Transaction protocol
5			
6			1 NB15_1.columns = NB15_features['Name'] 2 NB15_2.columns = NB15_features['Name'] 3 NB15_3.columns = NB15_features['Name'] 4 NB15_4.columns = NB15_features['Name'] 5 6 df_unsw = pd.concat([NB15_1, NB15_2, NB15_3, NB15_4], ignore_index=True)
7			
8	9	sttl	Integer Source to destination time to live value
9	10		1 df_unsw
10			
11	12	sloss	Integer Source packets retransmitted or dropped
12	13	srcip	srcip sport Integer dstip dsport proto state
13	dloss	Integer Destination packets retransmitted or dropped	
14	15	166.0.0.0	59.166.0.0 33661 149.171.126.9 1024 udp CON 0.036133 528 304 31 ... 0 2
15	Sload	1464	149.171.126.7 53 udp CON 0.001119 Source bits per second 146 178 31 ... 0 12
16	17	59.166.0.5	59.166.0.5 3593 149.171.126.5 53 udp CON 0.001209 132 164 31 ... 0 6
17	Spkts	49664	149.171.126.0 53 udp CON 0.001169 Source to destination packet count 146 178 31 ... 0 7
18	19	17	59.166.0.0 32119 149.171.126.9 111 udp CON 0.078339 568 312 31 ... 0 2
19	swin	integer	... Source TCP window advertisement value
20	2540038	59.166.0.5	59.166.0.5 33094 149.171.126.7 43433 tcp FIN 0.087306 320 1828 31 ... 1
21	2540039	59.166.0.7	59.166.0.7 20848 149.171.126.4 21 tcp CON 0.365058 Source TCP base sequence number 456 346 31 ... 2 2
22	2540040	59.166.0.3	59.166.0.3 21511 149.171.126.9 21 tcp CON 6.335154 1802 2088 31 ... 2 2
23	2540041	59.166.0.9	59.166.0.9 35433 149.171.126.0 80 tcp CON 2.200934 Mean of the ?ow packet size transmitted by the... 3498 166054 31 ... 1
24	2540042	175.45.176.0	175.45.176.0 17293 149.171.126.17 110 tcp CON 0.942984 574 676 62 ... 1
25	2540043	rows × 49 columns	
26	25	trans_depth	integer Represents the pipelined depth into the connec...
27	26	res_bdy_len	integer Actual uncompressed content size of the data t...
28			
29	27	Sjit	Float Source jitter (mSec)
30	28	Djit	Float Destination jitter (mSec)
31			
32	29	Time	Timestamp
33	30		record_start_time
34	31		
35	32		1 print("dataset shape: ",df_unsw.shape) 2 print(f"Memory usage: {df_unsw.memory_usage(deep=True).sum() / (1024**2):.2f} MB") 3 df_unsw.info()
36	33		
37	34		dataset shape: (2540043, 49) Float Memory usage: 1837.93 MB class pandas.core.frame.DataFrame RangeIndex: 2540043 entries, 0 to 2540042 Data columns (total 49 columns):
38	35		# Column Dtype --- srcip float64 Destination interpacket arrival time (mSec)
39	36		is_sm_ips_ports object If source (1) and destination (3)IP addresses ...
40	37		sport object No. for each state (6) according to specific r...
41	38		dstip object TCP connection setup time, the time between th...
42	39		ct_state_ttb integer dur float64 TCP connection setup round-trip time, the sum ...
43	40		shbytes int64 No. of flows that has methods such as Get and ...
44	41		dbbytes int64 If the ftp session is accessed by user and pas...
45	42		ct_ftp_cmd integer No. of flows that has a command in ftp session.
46	43		srtt int64 No. of connections that contain the same servi...
47	44		dtt1 int64 No. of connections that contain the same servi...
48	45		sloss int64 No. of connections of the same destination add...
49	46		dloss int64 No. of connections of the same source address ...
50	47		device int64 No. of connections of the same source address ...
51	48		Dload int64 No. of connections of the same source address ...
52	49		Spkts int64 No. of connections of the same source address (...)
53	50		Dpkts int64 No. of connections of the same source address (...)
54	51		swin int64 No. of connections of the same destination addr...
55	52		ct_dst_sport_ltm int64 No. of connections of the same desti...
56	53		ct_src_ip_ltm int64 No. of connections of the same source ip ...
57	54		ct_src_dport_ltm int64 No. of connections of the same desti...
58	55		ct_src_ip_ltm int64 No. of connections of the same source ip ...
59	56		ct_dst_ip_ltm int64 No. of connections of the same desti...
60	57		ct_src_dport_ltm int64 No. of connections of the same desti...
61	58		ct_dst_sport_ltm int64 No. of connections of the same desti...
62	59		ct_src_ip_ltm int64 No. of connections of the same source ip ...
63	60		ct_dst_ip_ltm int64 No. of connections of the same desti...
64	61		ct_src_dport_ltm int64 No. of connections of the same desti...
65	62		ct_dst_ip_ltm int64 No. of connections of the same desti...
66	63		ct_src_ip_ltm int64 No. of connections of the same source ip ...
67	64		ct_dst_dport_ltm int64 No. of connections of the same desti...
68	65		ct_src_ip_ltm int64 No. of connections of the same source ip ...
69	66		ct_dst_ip_ltm int64 No. of connections of the same desti...
70	67		ct_src_dport_ltm int64 No. of connections of the same desti...
71	68		ct_dst_ip_ltm int64 No. of connections of the same desti...
72	69		ct_src_ip_ltm int64 No. of connections of the same source ip ...
73	70		ct_dst_dport_ltm int64 No. of connections of the same desti...
74	71		ct_src_ip_ltm int64 No. of connections of the same source ip ...
75	72		ct_dst_ip_ltm int64 No. of connections of the same desti...
76	73		ct_src_dport_ltm int64 No. of connections of the same desti...
77	74		ct_dst_ip_ltm int64 No. of connections of the same desti...
78	75		ct_src_ip_ltm int64 No. of connections of the same source ip ...
79	76		ct_dst_dport_ltm int64 No. of connections of the same desti...
80	77		ct_src_ip_ltm int64 No. of connections of the same source ip ...
81	78		ct_dst_ip_ltm int64 No. of connections of the same desti...
82	79		ct_src_dport_ltm int64 No. of connections of the same desti...
83	80		ct_dst_ip_ltm int64 No. of connections of the same desti...
84	81		ct_src_ip_ltm int64 No. of connections of the same source ip ...
85	82		ct_dst_dport_ltm int64 No. of connections of the same desti...
86	83		ct_src_ip_ltm int64 No. of connections of the same source ip ...
87	84		ct_dst_ip_ltm int64 No. of connections of the same desti...
88	85		ct_src_dport_ltm int64 No. of connections of the same desti...
89	86		ct_dst_ip_ltm int64 No. of connections of the same desti...
90	87		ct_src_ip_ltm int64 No. of connections of the same source ip ...
91	88		ct_dst_dport_ltm int64 No. of connections of the same desti...
92	89		ct_src_ip_ltm int64 No. of connections of the same source ip ...
93	90		ct_dst_ip_ltm int64 No. of connections of the same desti...
94	91		ct_src_dport_ltm int64 No. of connections of the same desti...
95	92		ct_dst_ip_ltm int64 No. of connections of the same desti...
96	93		ct_src_ip_ltm int64 No. of connections of the same source ip ...
97	94		ct_dst_dport_ltm int64 No. of connections of the same desti...
98	95		ct_src_ip_ltm int64 No. of connections of the same source ip ...
99	96		ct_dst_ip_ltm int64 No. of connections of the same desti...
100	97		ct_src_dport_ltm int64 No. of connections of the same desti...
101	98		ct_dst_ip_ltm int64 No. of connections of the same desti...
102	99		ct_src_ip_ltm int64 No. of connections of the same source ip ...
103	100		ct_dst_dport_ltm int64 No. of connections of the same desti...
104	101		ct_src_ip_ltm int64 No. of connections of the same source ip ...
105	102		ct_dst_ip_ltm int64 No. of connections of the same desti...
106	103		ct_src_dport_ltm int64 No. of connections of the same desti...
107	104		ct_dst_ip_ltm int64 No. of connections of the same desti...
108	105		ct_src_ip_ltm int64 No. of connections of the same source ip ...
109	106		ct_dst_dport_ltm int64 No. of connections of the same desti...
110	107		ct_src_ip_ltm int64 No. of connections of the same source ip ...
111	108		ct_dst_ip_ltm int64 No. of connections of the same desti...
112	109		ct_src_dport_ltm int64 No. of connections of the same desti...
113	110		ct_dst_ip_ltm int64 No. of connections of the same desti...
114	111		ct_src_ip_ltm int64 No. of connections of the same source ip ...
115	112		ct_dst_dport_ltm int64 No. of connections of the same desti...
116	113		ct_src_ip_ltm int64 No. of connections of the same source ip ...
117	114		ct_dst_ip_ltm int64 No. of connections of the same desti...
118	115		ct_src_dport_ltm int64 No. of connections of the same desti...
119	116		ct_dst_ip_ltm int64 No. of connections of the same desti...
120	117		ct_src_ip_ltm int64 No. of connections of the same source ip ...
121	118		ct_dst_dport_ltm int64 No. of connections of the same desti...
122	119		ct_src_ip_ltm int64 No. of connections of the same source ip ...
123	120		ct_dst_ip_ltm int64 No. of connections of the same desti...
124	121		ct_src_dport_ltm int64 No. of connections of the same desti...
125	122		ct_dst_ip_ltm int64 No. of connections of the same desti...
126	123		ct_src_ip_ltm int64 No. of connections of the same source ip ...
127	124		ct_dst_dport_ltm int64 No. of connections of the same desti...
128	125		ct_src_ip_ltm int64 No. of connections of the same source ip ...
129	126		ct_dst_ip_ltm int64 No. of connections of the same desti...
130	127		ct_src_dport_ltm int64 No. of connections of the same desti...
131	128		ct_dst_ip_ltm int64 No. of connections of the same desti...
132	129		ct_src_ip_ltm int64 No. of connections of the same source ip ...
133	130		ct_dst_dport_ltm int64 No. of connections of the same desti...
134	131		ct_src_ip_ltm int64 No. of connections of the same source ip ...
135	132		ct_dst_ip_ltm int64 No. of connections of the same desti...
136	133		ct_src_dport_ltm int64 No. of connections of the same desti...
137	134		ct_dst_ip_ltm int64 No. of connections of the same desti...
138	135		ct_src_ip_ltm int64 No. of connections of the same source ip ...
139	136		ct_dst_dport_ltm int64 No. of connections of the same desti...
140	137		ct_src_ip_ltm int64 No. of connections of the same source ip ...
141	138		ct_dst_ip_ltm int64 No. of connections of the same desti...
142	139		ct_src_dport_ltm int64 No. of connections of the same desti...
143	140		ct_dst_ip_ltm int64 No. of connections of the same desti...
144	141		ct_src_ip_ltm int64 No. of connections of the same source ip ...
145	142		ct_dst_dport_ltm int64 No. of connections of the same desti...
146	143		ct_src_ip_ltm int64 No. of connections of the same source ip ...
147	144		ct_dst_ip_ltm int64 No. of connections of the same desti...
148	145		ct_src_dport_ltm int64 No. of connections of the same desti...
149	146		ct_dst_ip_ltm int64 No. of connections of the same desti...
150	147		ct_src_ip_ltm int64 No. of connections of the same source ip ...
151	148		ct_dst_dport_ltm int64 No. of connections of the same desti...
152	149		ct_src_ip_ltm int64 No. of connections of the same source ip ...
153	150		ct_dst_ip_ltm int64 No. of connections of the same desti...
154	151		ct_src_dport_ltm int64 No. of connections of the same desti...
155	152		ct_dst_ip_ltm int64 No. of connections of the same desti...
156	153		ct_src_ip_ltm int64 No. of connections of the same source ip ...
157	154		ct_dst_dport_ltm int64 No. of connections of the same desti...
158	155		ct_src_ip_ltm int64 No. of connections of the same source ip ...
159	156		ct_dst_ip_ltm int64 No. of connections of the same desti...
160	157		ct_src_dport_ltm int64 No. of connections of the same desti...
161	158		ct_dst_ip_ltm int64 No. of connections of the same desti...
162	159		ct_src_ip_ltm int64 No. of connections of the same source ip ...
163	160		ct_dst_dport_ltm int64 No. of connections of the same desti...
164	161		ct_src_ip_ltm int64 No. of connections of the same source ip ...
165	162		ct_dst_ip_ltm int64 No. of connections of the same desti...
166	163		ct_src_dport_ltm int64 No. of connections of the same desti...
167	164		ct_dst_ip_ltm int64 No. of connections of the same desti...
168	165		ct_src_ip_ltm int64 No. of connections of the same source ip ...
169	166		ct_dst_dport_ltm int64 No. of connections of the same desti...
170	167		ct_src_ip_ltm int64 No. of connections of the same source ip ...
171	168		ct_dst_ip_ltm int64 No. of connections of the same desti...
172	169		ct_src_dport_ltm int64 No. of connections of the same desti...
173	170		ct_dst_ip_ltm int64 No. of connections of the same desti...
174	171		ct_src_ip_ltm int64 No. of connections of the same source ip ...
175	172		ct_dst_dport_ltm int64 No. of connections of the same desti...
176	173		ct_src_ip_ltm int64 No. of connections of the same source ip ...
177	174		ct_dst_ip_ltm int64 No. of connections of the same desti...
178	175		ct_src_dport_ltm int64 No. of connections of the same desti...
179	176		ct_dst_ip_ltm int64 No. of connections of the same desti...
180	177		ct_src_ip_ltm int64 No. of connections of the same source ip ...
181	178		ct_dst_dport_ltm int64 No. of connections of the same desti...
182	179		ct_src_ip_ltm int64 No. of connections of the same source ip ...
183	180		ct_dst_ip_ltm int64 No. of connections of the same desti...
184	181		ct_src_dport_ltm int64 No. of connections of the same desti...
185	182		ct_dst_ip_ltm int64 No. of connections of the same desti...
186	183		ct_src_ip_ltm int64 No. of connections of the same source ip ...
187	184		ct_dst_dport_ltm int64 No. of connections of the same desti...
188	185		ct_src_ip_ltm int64 No. of connections of the same source ip ...
189	186		ct_dst_ip_ltm int64 No. of connections of the same desti...
190	187		ct_src_dport_ltm int64 No. of connections of the same desti...
191	188		ct_dst_ip_ltm int64 No. of connections of the same desti...
192	189		ct_src_ip_ltm int64 No. of connections of the same source ip ...
193	190		ct_dst_dport_ltm int64 No. of connections of the same desti...
194	191		ct_src_ip_ltm int64 No. of connections of the same source ip ...
195	192		ct_dst_ip_ltm int64 No. of connections of the same desti...
196	193		ct_src_dport_ltm int64 No. of connections of the same desti...
197	194		ct_dst_ip_ltm int64 No. of connections of the same desti...
198	195		ct_src_ip_ltm int64 No. of connections of the same source ip ...
199	196		ct_dst_dport_ltm int64 No. of connections of the same desti...
200	197		ct_src_ip_ltm int64 No. of connections of the same source ip ...
201	198		ct_dst_ip_ltm int64 No. of connections of the same desti...
202	199		ct_src_dport_ltm int64 No. of connections of the same desti...
203	200		ct_dst_ip_ltm int64 No. of connections of the same desti...
204	201		ct_src_ip_ltm int64 No. of connections of the same source ip ...
205	202		ct_dst_dport_ltm int64 No. of connections of the same desti...
206	203		ct_src_ip_ltm int64 No. of connections of the same source ip ...
207	204		ct_dst_ip_ltm int64 No. of connections of the same desti...
208	205		ct_src_dport_ltm int64 No. of connections of the same desti...
209	206		ct_dst_ip_ltm int64 No. of connections of the same desti...
210	207		ct_src_ip_ltm int64 No. of connections of the same source ip ...
211	208		ct_dst_dport_ltm int64 No. of connections of the same desti...
212	209		ct_src_ip_ltm int64 No. of connections of the same source ip ...
213	210		ct_dst_ip_ltm int64 No. of connections of the same desti...
214	211		ct_src_dport_ltm int64 No. of connections of the same desti...
215	212		ct_dst_ip_ltm int64 No. of connections of the same desti...
216	213		ct_src_ip_ltm int64 No. of connections of the same source ip ...
217	214		ct_dst_dport_ltm int64 No. of connections of the same desti...
218	215		ct_src_ip_ltm int64 No. of connections of the same source ip ...
219	216		ct_dst_ip_ltm int64 No. of connections of the same desti...
220	217		ct_src_dport_ltm int64 No. of connections of the same desti...
221	218		ct_dst_ip_ltm int64 No. of connections of the same desti...
222	219		ct_src_ip_ltm int64 No. of connections of the same source ip ...
223	220		ct_dst_dport_ltm int64 No. of connections of the same desti...
224	221		ct_src_ip_ltm int64 No. of connections of the same source ip ...
225	222		ct_dst_ip_ltm int64 No. of connections of the same desti...
226	223		ct_src_dport_ltm int64 No. of connections of the same desti...
227	224		ct_dst_ip_ltm int64 No. of connections of the same desti...
228	225		ct_src_ip_ltm int64 No. of connections of the same source ip ...
229	226		ct_dst_dport_ltm int64 No. of connections of the same desti...
230	227		ct_src_ip_ltm int64 No. of connections of the same source ip ...
231	228		ct_dst_ip_ltm int64 No. of connections of the same desti...
232	229		ct_src_d

```

19 dwin      int64
20 s47pb    ct_dst_src_ltmint64 integer  No of connections of the same source (1) and t...
21 dtcpb     int64
22 smeansz   attack_catint64 nominal   The name of each attack category. In this data...
23 dmeansz   int64
24 trans_depth Labelint64 binary       0 for normal and 1 for attack records
25 res_bdy_len int64
26 Sjit      float64
27 Djit      float64
28 Stime     int64
29 Ltime     int64
30 Sintpkt   float64
31 Dintpkt   float64
32 tcprtt    float64
33 synack    float64
34 ackdat    float64
35 is_sm_ips_ports int64
36 ct_state_ttl int64
37 ct_flw_http_mthd float64
38 is_ftp_login float64
39 ct_ftp_cmd   object
40 ct_srv_src   int64
41 ct_srv_dst   int64
42 ct_dst_ltm   int64
43 ct_src_ltm   int64
44 ct_src_dport_ltm int64
45 ct_dst_sport_ltm int64
46 ct_dst_src_ltm int64
47 attack_cat   object
48 Label      int64
dtypes: float64(12), int64(28), object(9)
memory usage: 949.6+ MB

```

Setup & work dir

```

1 import os, gc, math, json, numpy as np, pandas as pd
2 import matplotlib.pyplot as plt
3 from scipy import stats
4 from collections import Counter
5
6 BASE = "/content/drive/MyDrive/colab_zero_day_unsw"
7 EDA_DIR = f"{BASE}/eda"
8 os.makedirs(EDA_DIR, exist_ok=True)
9
10 def savefig(path):
11     plt.tight_layout()
12     plt.savefig(path, dpi=160, bbox_inches='tight')
13     plt.close()
14
15 def mem_mb(df):
16     return df.memory_usage(deep=True).sum()/1024**2
17
18 def numeric_columns(df, extra_drop=()):
19     return df.select_dtypes(include=[np.number]).columns.difference(list(extra_drop))

```

dataset cards & class/attack category balance

```

1 def dataset_card(name, df):
2     print(f"\n--- {name} ---")
3     print("shape:", df.shape, "| mem MB:", f"{mem_mb(df):.2f}")
4     if "Attack" in df.columns:
5         print("attack ratio:", df["Attack"].mean())
6     if "attack_cat" in df.columns:
7         print("attack_cat top 10:\n", df["attack_cat"].value_counts().head(10))
8
9 for nm, d in [("UNSW full", df_unsw),
10               ("Train (benign-only)", train_benign),
11               ("Validation (SEEN)", val_df),
12               ("Test REALISTIC (UNSEEN)", test_realistic),
13               ("Test STRESS (UNSEEN)", test_stress)]:
14     dataset_card(nm, d)
15
16 # Bars for attack categories (use full df attacks only to be representative)
17 if "attack_cat" in df_unsw.columns:

```

```

18     atk_counts = df_unsw[df_unsw["Label"]==1]["attack_cat"].value_counts().sort_values(ascending=False)
19     plt.figure(figsize=(8,4))
20     atk_counts.plot(kind="bar")
21     plt.title("UNSW-NB15 attack categories (full dataset)")
22     plt.ylabel("count"); plt.xlabel("attack_cat")
23     savefig(f"{EDA_DIR}/attack_cat_counts_full.png")
24
25 # Pies for splits
26 def pie_attack(df, title, fname):
27     if "Attack" not in df.columns: return
28     plt.figure(figsize=(4,4))
29     vals = df["Attack"].value_counts().reindex([0,1]).fillna(0)
30     plt.pie(vals, labels=["Normal","Attack"], autopct="%1.1f%%", startangle=90)
31     plt.title(title)
32     savefig(f"{EDA_DIR}/{fname}.png")
33
34 pie_attack(val_df, "Validation Attack vs Normal", "val_attack_pie")
35 pie_attack(test_realistic, "Test REALISTIC Attack vs Normal", "test_realistic_attack_pie")
36 pie_attack(test_stress, "Test STRESS Attack vs Normal", "test_stress_attack_pie")
37

```

==== UNSW full ====
shape: (2540043, 49) | mem MB: 1837.93
attack_cat top 10:
attack_cat
Generic 215481
Exploits 44525
Fuzzers 19195
DoS 16353
Reconnaissance 12228
Fuzzers 5051
Analysis 2677
Backdoor 1795
Reconnaissance 1759
Shellcode 1288
Name: count, dtype: int64

==== Train (benign-only) ====
shape: (2218760, 47) | mem MB: 285.67
attack ratio: 0.0
attack_cat top 10:
attack_cat
Fuzzers 0
Fuzzers 0
Reconnaissance 0
Shellcode 0
Analysis 0
Backdoor 0
Backdoors 0
DoS 0
Exploits 0
Generic 0
Name: count, dtype: int64

==== Validation (SEEN) ====
shape: (761893, 47) | mem MB: 98.11
attack ratio: 0.1263510755447287
attack_cat top 10:
attack_cat
Generic 64724
Exploits 13349
Fuzzers 5729
DoS 4923
Reconnaissance 3638
Fuzzers 1528
Analysis 792
Reconnaissance 532
Backdoor 518
Shellcode 376
Name: count, dtype: int64

==== Test REALISTIC (UNSEEN) ====
shape: (2219157, 47) | mem MB: 268.79
attack ratio: 0.00017889676124762692
attack_cat top 10:
attack_cat
Shellcode 223
Worms 174

missingness & basic numeric describe (memory-aware)

```

1 # choose a representative sample for heavy describe (to reduce RAM)
2 SAMPLE_N = min(len(df_unsw), 500_000)
3 df_sample = df_unsw.sample(n=SAMPLE_N, random_state=42) if len(df_unsw)>SAMPLE_N else df_unsw
4
5 # missingness on sample (full missingness can be huge to compute)
6 miss = df_sample.isna().sum().sort_values(ascending=False)
7 miss = miss[miss>0]
8 miss.to_csv(f"{EDA_DIR}/missingness_sample.csv")
9 print("Missing columns in sample (top):\n", miss.head(20))
10
11 # numeric describe on validation (stable for thresholding)
12 num_cols_val = numeric_columns(val_df, extra_drop=["Attack"])
13 desc = val_df[num_cols_val].describe(percentiles=[.01,.05,.25,.5,.75,.95,.99]).T
14 desc.to_csv(f"{EDA_DIR}/val_numeric_describe.csv")
15 print("Saved val numeric describe.")
16
Missing columns in sample (top):
Name
attack_cat      436655
is_ftp_login    281440
ct_flw_http_mthd 265283
dtype: int64
Saved val numeric describe.

```

✓ categoricals overview (proto, state, service, plus ports)

```

1 def bar_top(series, top=20, title="", fname=""):
2     vc = series.value_counts().head(top)
3     plt.figure(figsize=(8,4))
4     vc.plot(kind="bar")
5     plt.title(title); plt.ylabel("count")
6     savefig(f"{EDA_DIR}/{fname}.png")
7
8 for c in ["proto", "state", "service"]:
9     if c in df_unsw.columns:
10         bar_top(df_unsw[c], 20, f"{c} top-20 (full UNSW)", f"{c}_top20_full")
11
12 # Ports: ensure numeric
13 for pcol in ["sport", "dsport"]:
14     if pcol in df_unsw.columns and not np.issubdtype(df_unsw[pcol].dtype, np.number):
15         df_unsw[pcol] = pd.to_numeric(df_unsw[pcol], errors="coerce")
16
17 for pcol in ["sport", "dsport"]:
18     if pcol in df_unsw.columns:
19         bar_top(df_unsw[pcol], 20, f"{pcol} top-20 (full UNSW)", f"{pcol}_top20_full")
20

```

✓ feature separation (KS & Cohen's d) on validation set

```

1 def feature_separation(df, cols, target="Attack"):
2     rows=[]
3     a = df[df[target]==1][cols]
4     n = df[df[target]==0][cols]
5     for c in cols:
6         try:
7             ks_stat, ks_p = stats.ks_2samp(a[c].values, n[c].values, alternative='two-sided', mode='auto')
8             mu1, mu0 = a[c].mean(), n[c].mean()
9             s1, s0 = a[c].std(ddof=1), n[c].std(ddof=1)
10            n1, n0 = a[c].shape[0], n[c].shape[0]
11            sp = math.sqrt(((n1-1)*s1**2 + (n0-1)*s0**2) / max(n1+n0-2,1))
12            d = (mu1 - mu0) / (sp + 1e-12)
13            rows.append((c, ks_stat, ks_p, d, mu0, mu1))
14        except Exception as e:
15            rows.append((c, np.nan, np.nan, np.nan, np.nan, np.nan))
16    out = pd.DataFrame(rows, columns=["feature", "KS", "KS_p", "Cohen_d", "mean_normal", "mean_attack"])
17    out["abs_d"] = out["Cohen_d"].abs()
18    out.sort_values(["KS", "abs_d"], ascending=[False, False], inplace=True)

```

```

19     return out
20
21 num_cols_val = numeric_columns(val_df, extra_drop=["Attack"])
22 sep_tbl = feature_separation(val_df, num_cols_val)
23 sep_tbl.to_csv(f"{EDA_DIR}/feature_separation_val.csv", index=False)
24 print("Top 10 separating features:\n", sep_tbl.head(10))
25

Top 10 separating features:
      feature    KS  KS_p      Cohen_d  mean_normal  mean_attack \
4       Label  1.000000  0.0  1.000000e+12  0.000000e+00   1.000000
37      sttl  0.978912  0.0  6.397585e+00  3.706996e+01  240.080194
19  ct_state_ttl  0.977707  0.0  5.442588e+00  3.365999e-02   1.830241
2       Dload  0.801384  0.0 -6.769236e-01  2.799810e+06  11541.887695
22     dmeansz  0.774339  0.0 -8.500083e-01  3.114549e+02   36.684364
20     dbytes  0.708067  0.0 -2.208217e-01  4.124267e+04  5080.558640
36   state_code  0.707914  0.0  1.076418e+00  4.363343e+00   5.811574
3       Dpkts  0.707901  0.0 -3.409656e-01  4.821069e+01   6.275456
0       Dintpkt  0.707664  0.0 -3.172701e-02  8.410172e+01  38.834361
24      dttl  0.705257  0.0  4.061654e-01  2.854549e+01  45.753890

      abs_d
4  1.000000e+12
37 6.397585e+00
19 5.442588e+00
2 6.769236e-01
22 8.500083e-01
20 2.208217e-01
36 1.076418e+00
3 3.409656e-01
0 3.172701e-02
24 4.061654e-01

```

✓ histograms for top 12 features (validation, clipped 1–99%)

```

1 TOPK = 12
2 top_feats = sep_tbl["feature"].head(TOPK).tolist()
3
4 def plot_hist_by_class(df, cols, title_prefix, fname_prefix):
5     for c in cols:
6         plt.figure(figsize=(5,3))
7         x0 = df[df["Attack"]==0][c].values
8         x1 = df[df["Attack"]==1][c].values
9         lo = np.nanpercentile(df[c].values, 1)
10        hi = np.nanpercentile(df[c].values, 99)
11        bins = 50
12        plt.hist(np.clip(x0, lo, hi), bins=bins, alpha=0.6, label="Normal", density=True)
13        plt.hist(np.clip(x1, lo, hi), bins=bins, alpha=0.6, label="Attack", density=True)
14        plt.xlabel(c); plt.ylabel("density")
15        plt.title(f"{title_prefix}: {c}")
16        plt.legend()
17        savefig(f"{EDA_DIR}/{fname_prefix}_{c}.png")
18
19 plot_hist_by_class(val_df, top_feats, "Validation distributions", "val_hist")
20

```

✓ correlation heatmap (numeric, on a sample to save RAM) + high-corr pairs

```

1 # sample from validation for correlation
2 VAL_CORR_N = min(len(val_df), 150_000)
3 val_samp = val_df.sample(n=VAL_CORR_N, random_state=42) if len(val_df)>VAL_CORR_N else val_df
4 num_cols_val = numeric_columns(val_samp, extra_drop=["Attack"])
5 corr = val_samp[num_cols_val].corr().fillna(0.0)
6
7 plt.figure(figsize=(8,6))
8 plt.imshow(corr.values, aspect='auto', interpolation='nearest')
9 plt.colorbar(label="Pearson r")
10 plt.title("Validation correlation heatmap (sampled)")
11 plt.xticks([], []); plt.yticks([], []) # hide tick clutter for large matrices
12 savefig(f"{EDA_DIR}/corr_heatmap_val_sampled.png")
13
14 pairs = []
15 thr = 0.98

```

```

16 cols = list(num_cols_val)
17 for i in range(len(cols)):
18     for j in range(i+1, len(cols)):
19         r = corr.iat[i,j]
20         if abs(r) >= thr:
21             pairs.append((cols[i], cols[j], float(r)))
22 pairs = sorted(pairs, key=lambda x: -abs(x[2]))
23 pd.DataFrame(pairs, columns=["feat_a","feat_b","r"]).to_csv(f"{EDA_DIR}/high_corr_pairs.csv", index=False)
24 print("High-corr pairs (|r|>=0.98):", len(pairs))
25

```

High-corr pairs (|r|>=0.98): 3

✓ outlier/tail heaviness (validation)

```

1 def tail_heaviness(df, cols):
2     rows=[]
3     for c in cols:
4         x = df[c].values
5         q1,q3 = np.nanpercentile(x, [25,75]); iqr = q3 - q1
6         p01,p99 = np.nanpercentile(x, [1,99])
7         kurt = stats.kurtosis(x, fisher=True, nan_policy='omit')
8         skw = stats.skew(x, nan_policy='omit')
9         out_frac = np.mean((x < q1 - 3*iqr) | (x > q3 + 3*iqr))
10        rows.append((c, skw, kurt, out_frac, p01, p99, q1, q3))
11    T = pd.DataFrame(rows, columns=["feature", "skew", "kurtosis", "outlier_frac", "p01", "p99", "q1", "q3"])
12    return T.sort_values("outlier_frac", ascending=False)
13
14 tails = tail_heaviness(val_df, numeric_columns(val_df, extra_drop=["Attack"]))
15 tails.to_csv(f"{EDA_DIR}/tail_heaviness_val.csv", index=False)
16
17 plt.figure(figsize=(8,4))
18 top_out = tails.head(20)
19 plt.bar(range(len(top_out)), top_out["outlier_frac"])
20 plt.xticks(range(len(top_out)), top_out["feature"], rotation=90, fontsize=7)
21 plt.ylabel("fraction beyond 3*IQR")
22 plt.title("Top-20 heavy-tail/outlier features (validation)")
23 savefig(f"{EDA_DIR}/outlier_frac_bar.png")
24
25 heavy_tail_frac = (tails["outlier_frac"] > 0.01).mean()
26 suggest_scaler = "RobustScaler" if heavy_tail_frac > 0.1 else "StandardScaler"
27 print(f"Suggested scaler: {suggest_scaler} (heavy-tail features: {heavy_tail_frac*100:.1f}% )")
28

```

/usr/local/lib/python3.12/dist-packages/scipy/stats/_stats_py.py:1231: RuntimeWarning: overflow encountered in square
s = s**2
/usr/local/lib/python3.12/dist-packages/numpy/_core/_methods.py:127: RuntimeWarning: overflow encountered in reduce
ret = umr_sum(arr, axis, dtype, out, keepdims, where=where)
Suggested scaler: RobustScaler (heavy-tail features: 88.1%)

✓ time profile (if Stime exists) + per-category rates

```

1 if "Stime" in df_unsw.columns:
2     # Stime is epoch (seconds). Plot counts by hour on a sample to save RAM.
3     S_N = min(len(df_unsw), 500_000)
4     s = df_unsw.sample(n=S_N, random_state=42) if len(df_unsw)>S_N else df_unsw
5     ts = pd.to_datetime(s["Stime"], unit="s", errors="coerce")
6     hours = ts.dt.floor('H')
7     counts = hours.value_counts().sort_index()
8     plt.figure(figsize=(9,3))
9     counts.plot()
10    plt.title("Traffic volume over time (sampled, hourly)")
11    plt.ylabel("flows/hour"); plt.xlabel("time")
12    savefig(f"{EDA_DIR}/time_series_hourly.png")
13
14 # Attack rate per attack_cat (on full df)
15 if "attack_cat" in df_unsw.columns and "Label" in df_unsw.columns:
16     grp = df_unsw.groupby("attack_cat")["Label"].agg(['count','mean']).sort_values('count', ascending=False)
17     grp.rename(columns={'mean':'attack_ratio'}, inplace=True)
18     grp.to_csv(f"{EDA_DIR}/attack_cat_stats.csv")
19     print("Saved attack_cat_stats.csv")
20

/tmppython-input-1388815498.py:6: FutureWarning: 'H' is deprecated and will be removed in a future version, please use 'h' instead
hours = ts.dt.floor('H')
Saved attack_cat_stats.csv

```

EDA-driven preprocessing plan

```

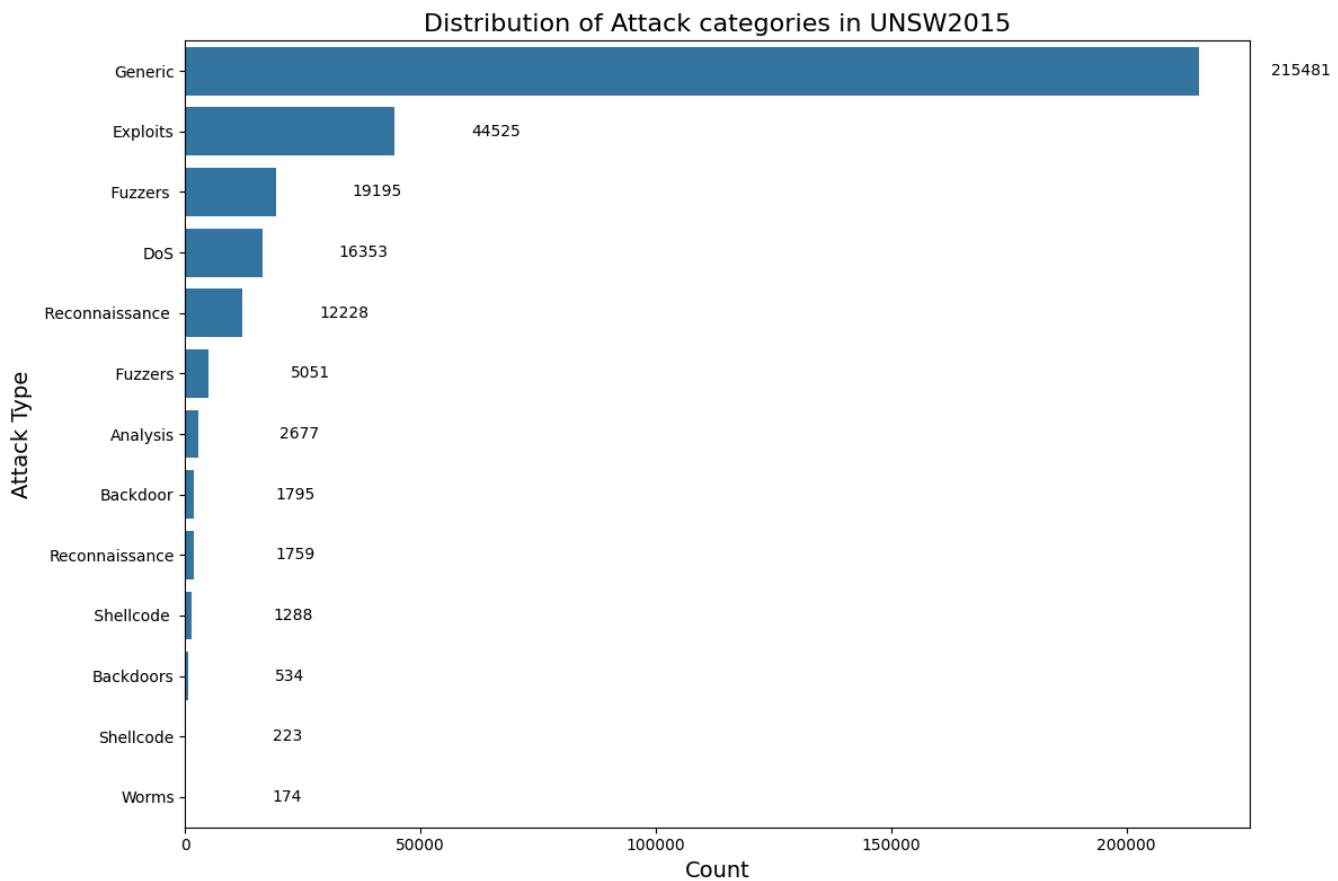
1 # drop ID/time/seq columns; keep ports and proto/state/service; use RobustScaler if many heavy tails
2 drop_cols_proposed = [c for c in ["srcip","dstip","Stime","Ltime","stcpb","dtcpb"] if c in df_unsw.columns]
3 cat_cols_proposed = [c for c in ["proto","state","service"] if c in df_unsw.columns]
4
5 num_cols_proposed = df_unsw.select_dtypes(include=[np.number]).columns.tolist()
6 num_cols_proposed = [c for c in num_cols_proposed if c not in set(["Label","Attack"])]
7
8 preproc_plan_unsw = {
9     "drop_cols": drop_cols_proposed,
10    "categoricals": cat_cols_proposed,
11    "numeric_candidates": num_cols_proposed,
12    "scaler": suggest_scaler,
13    "top_sep_features_val": sep_tbl["feature"].head(20).tolist(),
14    "high_corr_pairs_thresh": 0.98
15 }
16 with open(f"{EDA_DIR}/preprocessing_plan_unsw.json","w") as f:
17     json.dump(preproc_plan_unsw, f, indent=2)
18
19 print("Saved preprocessing_plan_unsw.json")
20

```

```

1 import matplotlib.pyplot as plt
2 import seaborn as sns
3
4 target_col = 'attack_cat'
5 plt.figure(figsize=(12, 8))
6 ax = sns.countplot(df_unsw, y=target_col, order=df_unsw[target_col].value_counts().index)
7
8 # 3. Add annotations and styling
9 plt.title('Distribution of Attack categories in UNSW2015', fontsize=16)
10 plt.xlabel('Count', fontsize=14)
11 plt.ylabel('Attack Type', fontsize=14)
12
13 # Add count values on each bar
14 for p in ax.patches:
15     width = p.get_width()
16     plt.text(width + 0.1 * max(df_unsw[target_col].value_counts()), 
17               p.get_y() + p.get_height()/2.,
18               f'{int(width)}',
19               ha='center', va='center')
20
21 plt.tight_layout()
22 plt.show()

```



```

1 # 1.1 Fix odd column names (spaces, case)
2 df = df.rename(columns=lambda c: c.strip().replace(" ", "_"))
3 # e.g., 'ct_src_ltm' -> 'ct_src_ltm' above; normalize further:
4 df = df.rename(columns={"ct_src_ltm": "ct_src_ltm"})
5
6 # 1.2 Coerce ports to numeric (often object in raw CSVs)
7 for col in ["sport", "dsport", "ct_ftp_cmd"]:
8     if col in df.columns and df[col].dtype == "object":
9         df[col] = pd.to_numeric(df[col], errors="coerce")
10
11 # 1.3 Convert low-cardinality text to category (saves RAM; fine for OneHot later)
12 for col in ["proto", "state", "service", "srcip", "dstip", "attack_cat"]:
13     if col in df.columns and df[col].dtype == "object":
14         # IPs are high-card; still set to category for cheaper memory even if we later drop them
15         df[col] = df[col].astype("category")
16
17 # 1.4 Downcast numerics where safe
18 def downcast_dfx():
19     for c in dfx.select_dtypes(include=["int64"]).columns:
20         dfx[c] = pd.to_numeric(dfx[c], downcast="unsigned") # counts/IDs are nonnegative in UNSW
21     for c in dfx.select_dtypes(include=["float64"]).columns:
22         dfx[c] = pd.to_numeric(dfx[c], downcast="float")
23     return dfx
24
25 df = downcast_dfx()
26 print(df.dtypes.value_counts())
27 print("RAM (MB):", df.memory_usage(deep=True).sum()/1024**2)
28

```

uint8	15
float32	11
uint32	7
uint16	6
float64	4
category	1
category	1
category	1

```
category    1
category    1
category    1
Name: count, dtype: int64
RAM (MB): 334.309627532959
```

```
1 # 2.1 Normalize labels
2 assert "Label" in df.columns and "attack_cat" in df.columns, "Expected Label & attack_cat columns."
3 df["Attack"] = df["Label"].astype("int8") # 1=attack, 0=normal
4
5 # 2.2 Pick UNSEEN zero-day categories (least frequent TWO, reproducibly)
6 atk_only = df[df["Attack"]==1]
7 vc = atk_only["attack_cat"].value_counts()
8 # Keep only valid strings
9 vc = vc[vc.index.notna()]
10 unseen_cats = list(vc.sort_values(ascending=True).head(2).index) # e.g., ['Worms', 'Shellcode'] usually
11 print("Zero-day UNSEEN categories:", unseen_cats)
12
13 is_unseen = df["attack_cat"].isin(unseen_cats)
14 df_seen = df[~is_unseen] | (df["Attack"]==0)].copy() # all benign + seen attacks
15 df_unseen = df[is_unseen] | (df["Attack"]==0)].copy() # all benign + unseen attacks
16
17 print("SEEN:", df_seen.shape, "UNSEEN:", df_unseen.shape)
18
```

```
Zero-day UNSEEN categories: ['Worms', 'Shellcode']
SEEN: (2539646, 50) UNSEEN: (2219157, 50)
```

```
1 from sklearn.model_selection import train_test_split
2
3 # 3.1 Train on benign-only from SEEN
4 train_benign = df_seen[df_seen["Attack"]==0].copy()
5
6 # 3.2 Validation (SEEN): stratified sample for stable thresholding
7 val_frac = 0.30
8 val_df, _ = train_test_split(
9     df_seen, test_size=(1 - val_frac),
10    stratify=df_seen["Attack"], random_state=42
11 )
12
13 # 3.3 Test sets from UNSEEN
14 test_df = df_unseen.copy()
15
16 # REALISTIC prevalence (e.g., 10% attacks)
17 target_prev = 0.10
18 ben = test_df[test_df["Attack"]==0]
19 atk = test_df[test_df["Attack"]==1]
20 n_b = len(ben)
21 n_a = min(len(atk), int(n_b * target_prev / (1 - target_prev)))
22 atk_down = atk.sample(n=n_a, random_state=42) if n_a>0 else atk
23 test_realistic = pd.concat([ben, atk_down]).sample(frac=1.0, random_state=42).reset_index(drop=True)
24
25 # STRESS: as-is
26 test_stress = test_df.sample(frac=1.0, random_state=42).reset_index(drop=True)
27
28 for nm, d in [("train_benign", train_benign), ("val_df", val_df),
29                 ("test_realistic", test_realistic), ("test_stress", test_stress)]:
30     print(nm, d.shape, "attack ratio:", d["Attack"].mean())
31
32 # Save splits
33 split_dir = f"{BASE}/splits"
34 train_benign.to_parquet(f"{split_dir}/train_benign_seen.parquet", index=False)
35 val_df.to_parquet(f"{split_dir}/val_seen.parquet", index=False)
36 test_realistic.to_parquet(f"{split_dir}/test_unseen_realistic.parquet", index=False)
37 test_stress.to_parquet(f"{split_dir}/test_unseen_stress.parquet", index=False)
38 with open(f"{split_dir}/meta.json", "w") as f:
39     json.dump({"unseen_categories": [str(x) for x in unseen_cats],
40                "target_attack_prevalence_realistic": target_prev}, f, indent=2)
41
```

```
train_benign (2218760, 50) attack ratio: 0.0
val_df (761893, 50) attack ratio: 0.1263510755447287
test_realistic (2219157, 50) attack ratio: 0.00017889676124762692
test_stress (2219157, 50) attack ratio: 0.00017889676124762692
```

```

1 drop_cols = [c for c in ["srcip","dstip","Stime","Ltime","stcpb","dtcpb"] if c in df.columns]
2 cat_cols = [c for c in ["proto","state","service"] if c in df.columns]
3 label_cols = ["Label","attack_cat","Attack"]
4
5 # numeric candidates = all numeric minus drops & labels
6 num_cols = df.select_dtypes(include=["float32","float64","int16","int32","int64","uint8","uint16","uint32"]).columns.tolist()
7 num_cols = [c for c in num_cols if c not in set(drop_cols + label_cols)]
8
9 # remove ports if you choose not to use them
10 num_cols = [c for c in num_cols if c not in ("sport","dsport")]
11
12 print("Drop cols:", drop_cols)
13 print("Categoricals:", cat_cols)
14 print("Numeric count:", len(num_cols))
15

```

Drop cols: ['srcip', 'dstip', 'Stime', 'Ltime', 'stcpb', 'dtcpb']
 Categoricals: ['proto', 'state', 'service']
 Numeric count: 38

```

1 import numpy as np, pandas as pd, gc
2
3 # Categoricals to keep (if present)
4 cat_cols = [c for c in ["proto","state","service"] if c in train_benign.columns]
5 label_cols = ["Label","attack_cat","Attack"]
6
7 # Build code maps on TRAIN BENIGN ONLY (no leakage)
8 code_maps = {}
9 for c in cat_cols:
10     cats = pd.Categorical(train_benign[c]).categories.tolist()
11     code_maps[c] = {v:i for i,v in enumerate(cats)}
12
13 # FIXED encoder: map from object, allow -1 for unseen categories
14 def add_codes(df_in):
15     df_out = df_in # in-place to save RAM
16     for c in cat_cols:
17         cmap = code_maps[c]
18         df_out[c + "_code"] = df_out[c].astype("object").map(cmap).fillna(-1).astype("int16")
19     return df_out
20
21 # Apply to all splits
22 train_benign = add_codes(train_benign)
23 val_df      = add_codes(val_df)
24 test_realistic = add_codes(test_realistic)
25 test_stress   = add_codes(test_stress)
26
27 code_cols = [c+"_code" for c in cat_cols]
28
29 # Drop high-card/ID/leakage columns if present
30 drop_cols = [c for c in ["srcip","dstip","Stime","Ltime","stcpb","dtcpb"] if c in train_benign.columns]
31 for df_ in (train_benign, val_df, test_realistic, test_stress):
32     for c in drop_cols:
33         if c in df_.columns: df_.drop(columns=c, inplace=True)
34
35 # Numeric base features (we add code_cols later)
36 num_base_cols = train_benign.select_dtypes(include=[np.number]).columns.tolist()
37 num_base_cols = [c for c in num_base_cols if c not in set(label_cols + code_cols)]
38 print("num_base_cols:", len(num_base_cols), " | code_cols:", code_cols)
39 gc.collect();
40

```

num_base_cols: 38 | code_cols: ['proto_code', 'state_code', 'service_code']

```

1 def compute_clip_bounds(df_num, lower=1.0, upper=99.0):
2     lo = df_num.quantile(lower/100.0)
3     hi = df_num.quantile(upper/100.0)
4     return lo.to_dict(), hi.to_dict()
5
6 train_benign_num = train_benign[num_base_cols].astype("float32", copy=False)
7 lo_b, hi_b = compute_clip_bounds(train_benign_num, 1.0, 99.0)
8 del train_benign_num; gc.collect()
9
10 def clip_numeric_block(df_in):
11     Xn = df_in[num_base_cols].astype("float32", copy=True)
12     for c in num_base_cols:
13         Xn[c] = np.clip(Xn[c].values, lo_b[c], hi_b[c])

```

```

14     return Xn
15
16 def build_matrix(df_in):
17     # numeric (winsorized) + categorical codes (as float32)
18     Xn = clip_numeric_block(df_in)
19     Xc = df_in[code_cols].astype("float32", copy=False) if code_cols else None
20     return Xn if Xc is None else pd.concat([Xn, Xc], axis=1)
21

```

```

1 from sklearn.preprocessing import RobustScaler, StandardScaler
2
3 # Sample benign for fitting scaler (keeps RAM/time low)
4 N_TRAIN_MAX = 400_000
5 rs = np.random.RandomState(42)
6 idx_fit = rs.choice(len(train_benign), size=min(N_TRAIN_MAX, len(train_benign)), replace=False)
7 train_sample = train_benign.iloc[idx_fit]
8
9 Scaler = RobustScaler # robust to heavy tails
10 scaler = Scaler(with_centering=True, with_scaling=True)
11
12 X_train_sample_df = build_matrix(train_sample)
13 scaler.fit(X_train_sample_df.values)
14 del X_train_sample_df, train_sample; gc.collect()
15
16 def transform_split(df_in):
17     Xdf = build_matrix(df_in)
18     X = scaler.transform(Xdf.values).astype("float32", copy=False)
19     del Xdf; gc.collect()
20     return X
21
22 X_train = transform_split(train_benign)      # benign-only (you can subsample for fit below)
23 X_val   = transform_split(val_df)
24 X_tr    = transform_split(test_realistic)
25 X_ts    = transform_split(test_stress)
26
27 y_val = val_df["Attack"].astype(int).values
28 y_tr  = test_realistic["Attack"].astype(int).values
29 y_ts  = test_stress["Attack"].astype(int).values
30
31 print("Shapes -> X_train:", X_train.shape, " X_val:", X_val.shape, " X_tr:", X_tr.shape, " X_ts:", X_ts.shape)
32 gc.collect();
33

```

Shapes -> X_train: (2218760, 41) X_val: (761893, 41) X_tr: (2219157, 41) X_ts: (2219157, 41)

```

1 N_FIT_MAX = 400_000
2 if len(X_train) > N_FIT_MAX:
3     idx_fit2 = rs.choice(len(X_train), size=N_FIT_MAX, replace=False)
4     X_fit = X_train[idx_fit2]
5 else:
6     X_fit = X_train
7 print("Fitting on:", X_fit.shape)
8

```

Fitting on: (400000, 41)

```

1 from sklearn.metrics import precision_recall_curve, classification_report, roc_auc_score, confusion_matrix, roc_curve, precision_score
2 import numpy as np
3
4 def pick_thresh_by_Fbeta(y_true, scores, beta=2.0):
5     P, R, thr = precision_recall_curve(y_true, scores)
6     P, R = P[:-1], R[:-1]
7     fbeta = (1+beta**2) * (P*R) / (beta**2 * P + R + 1e-12)
8     i = np.nanargmax(fbeta)
9     return float(thr[i]), float(P[i]), float(R[i]), float(fbeta[i])
10
11 def pick_thresh_at_fpr(y_true, scores, max_fpr=0.05):
12     fpr, tpr, thr = roc_curve(y_true, scores)
13     ok = np.where(fpr <= max_fpr)[0]
14     if len(ok)==0:
15         j = np.argmax(tpr - fpr)
16         return float(thr[j]), float(fpr[j]), float(tpr[j])
17     i = ok[np.argmax(tpr[ok])]
18     return float(thr[i]), float(fpr[i]), float(tpr[i])
19

```

```

20 def summarize(y_true, y_pred, scores, name=""):
21     P,R,F1,_ = precision_recall_fscore_support(y_true, y_pred, average='binary', zero_division=0)
22     auc_ = roc_auc_score(y_true, scores) if len(np.unique(y_true))>1 else np.nan
23     cm = confusion_matrix(y_true, y_pred)
24     print(f"\n== {name} ==")
25     print(f"Precision={P:.4f} Recall={R:.4f} F1={F1:.4f} ROC-AUC={auc_:.4f}")
26     print("Confusion matrix:\n", cm)
27     print(classification_report(y_true, y_pred, digits=4))
28

```

Train and Evaluate Isolation Forest

```

1 from sklearn.ensemble import IsolationForest
2
3 iforest = IsolationForest(
4     n_estimators=400, max_samples='auto', contamination='auto',
5     random_state=42, n_jobs=-1
6 ).fit(X_fit)
7
8 val_scores = -iforest.score_samples(X_val)
9 tr_scores = -iforest.score_samples(X_tr)
10 ts_scores = -iforest.score_samples(X_ts)
11
12 # F2-optimal threshold (recall-oriented)
13 th_f2, p_f2, r_f2, f2 = pick_thresh_by_Fbeta(y_val, val_scores, beta=2.0)
14 y_tr_f2 = (tr_scores >= th_f2).astype(int)
15 y_ts_f2 = (ts_scores >= th_f2).astype(int)
16 print(f"IForest F2 threshold: {th_f2:.6f} (val P={p_f2:.3f}, R={r_f2:.3f}, F2={f2:.3f})")
17 summarize(y_tr, y_tr_f2, tr_scores, "IForest | REALISTIC | F2-opt")
18 summarize(y_ts, y_ts_f2, ts_scores, "IForest | STRESS | F2-opt")
19
20 # FPR-controlled threshold (<5% on validation)
21 th_fpr, fpr_sel, tpr_sel = pick_thresh_at_fpr(y_val, val_scores, max_fpr=0.05)
22 y_tr_fpr = (tr_scores >= th_fpr).astype(int)
23 y_ts_fpr = (ts_scores >= th_fpr).astype(int)
24 print(f"IForest FPR<=5% threshold: {th_fpr:.6f} (val FPR={fpr_sel:.3f}, TPR={tpr_sel:.3f})")
25 summarize(y_tr, y_tr_fpr, tr_scores, "IForest | REALISTIC | FPR<=5%")
26 summarize(y_ts, y_ts_fpr, ts_scores, "IForest | STRESS | FPR<=5%")
27

```

IForest F2 threshold: 0.505144 (val P=0.578, R=0.993, F2=0.868)

== IForest | REALISTIC | F2-opt ==
Precision=0.0017 Recall=0.9975 F1=0.0034 ROC-AUC=0.9615
Confusion matrix:
[[1986445 232315]
[1 396]]
precision recall f1-score support
0 1.0000 0.8953 0.9448 2218760
1 0.0017 0.9975 0.0034 397
accuracy 0.8953 2219157
macro avg 0.5009 0.9464 0.4741 2219157
weighted avg 0.9998 0.8953 0.9446 2219157

== IForest | STRESS | F2-opt ==
Precision=0.0017 Recall=0.9975 F1=0.0034 ROC-AUC=0.9615
Confusion matrix:
[[1986445 232315]
[1 396]]
precision recall f1-score support
0 1.0000 0.8953 0.9448 2218760
1 0.0017 0.9975 0.0034 397
accuracy 0.8953 2219157
macro avg 0.5009 0.9464 0.4741 2219157
weighted avg 0.9998 0.8953 0.9446 2219157

IForest FPR<=5% threshold: 0.551726 (val FPR=0.050, TPR=0.479)

== IForest | REALISTIC | FPR<=5% ==
Precision=0.0022 Recall=0.6196 F1=0.0044 ROC-AUC=0.9615
Confusion matrix:
[[2108237 110523]]

```
[ 151  246]]
precision    recall   f1-score   support
          0    0.9999    0.9502    0.9744    2218760
          1    0.0022    0.6196    0.0044     397

accuracy                           0.9501    2219157
macro avg                           0.5011    0.7849    0.4894    2219157
weighted avg                          0.9997    0.9501    0.9742    2219157

== IForest | STRESS      | FPR<=5% ==
Precision=0.0022 Recall=0.6196 F1=0.0044 ROC-AUC=0.9615
Confusion matrix:
[[2108237  110523]
 [ 151    246]]
precision    recall   f1-score   support
          0    0.9999    0.9502    0.9744    2218760
          1    0.0022    0.6196    0.0044     397
```

```
1 import numpy as np
2
3 def check_finiteness(X, name):
4     n = X.shape[0]
5     bad_nan = np.isnan(X).sum()
6     bad_inf = np.isinf(X).sum()
7     print(f"{name}: NaN={bad_nan}, Inf={bad_inf}, shape={X.shape}")
8
9 check_finiteness(X_train, "X_train")
10 check_finiteness(X_val, "X_val")
11 check_finiteness(X_tr, "X_tr")
12 check_finiteness(X_ts, "X_ts")
13

X_train: NaN=3330904, Inf=0, shape=(2218760, 41)
X_val: NaN=1263517, Inf=0, shape=(761893, 41)
X_tr: NaN=3331234, Inf=0, shape=(2219157, 41)
X_ts: NaN=3331234, Inf=0, shape=(2219157, 41)
```

```
1 import numpy as np, pandas as pd, gc
2 from sklearn.preprocessing import RobustScaler
3
4 # 1) Identify the columns we intend to use
5 label_cols = ["Label", "attack_cat", "Attack"]
6 code_cols = [c for c in ["proto_code", "state_code", "service_code"] if c in train_benign.columns]
7
8 # Base numeric candidates = all numeric except labels & codes
9 num_base_cols = train_benign.select_dtypes(include=[np.number]).columns.tolist()
10 num_base_cols = [c for c in num_base_cols if c not in set(label_cols + code_cols)]
11
12 # 2) Drop columns that are ALL-NaN in train_benign
13 all_nan_cols = train_benign[num_base_cols].isna().all(axis=0)
14 drop_all_nan = all_nan_cols[all_nan_cols].index.tolist()
15 if drop_all_nan:
16     print("Dropping all-NaN numeric cols:", drop_all_nan)
17 num_base_cols = [c for c in num_base_cols if c not in set(drop_all_nan)]
18
19 # 3) Compute clip bounds (1st/99th) on train_benign and FIX any NaN bounds by using the median
20 def compute_clip_bounds_safe(df_num, lower=1.0, upper=99.0):
21     lo = df_num.quantile(lower/100.0)
22     hi = df_num.quantile(upper/100.0)
23     med = df_num.median()
24     # Replace NaN lo/hi with median; if still NaN, use 0.0
25     lo = lo.fillna(med).fillna(0.0)
26     hi = hi.fillna(med).fillna(0.0)
27     # Ensure lo <= hi
28     bad = lo > hi
29     if bad.any():
30         # when inverted, set both to median
31         lo[bad] = med[bad]
32         hi[bad] = med[bad]
33     return lo.to_dict(), hi.to_dict()
34
35 train_benign_num = train_benign[num_base_cols].astype("float32", copy=False)
36 lo_b, hi_b = compute_clip_bounds_safe(train_benign_num, 1.0, 99.0)
37 del train_benign_num; gc.collect()
```

```

39 # 4) Clip numerics and build matrices (numeric + code cols)
40 def clip_numeric_block(df_in, cols, lo_b, hi_b):
41     Xn = df_in[cols].astype("float32", copy=True)
42     for c in cols:
43         lo, hi = lo_b[c], hi_b[c]
44         # safe clip; if lo==hi, this becomes a constant column (we'll drop constants next)
45         Xn[c] = np.clip(Xn[c].values, lo, hi)
46     return Xn
47
48 def build_matrix(df_in):
49     Xn = clip_numeric_block(df_in, num_base_cols, lo_b, hi_b)
50     Xc = df_in[code_cols].astype("float32", copy=False) if code_cols else None
51     X = Xn if Xc is None else pd.concat([Xn, Xc], axis=1)
52     return X
53
54 X_train_df = build_matrix(train_benign)
55 X_val_df   = build_matrix(val_df)
56 X_tr_df    = build_matrix(test_realistic)
57 X_ts_df    = build_matrix(test_stress)
58
59 # 5) Drop constant columns (zero variance on TRAIN)
60 stds = X_train_df.std(axis=0, ddof=0)
61 const_cols = stds[~np.isfinite(stds) | (stds == 0)].index.tolist()
62 if const_cols:
63     print("Dropping constant/invalid cols:", const_cols)
64     X_train_df.drop(columns=const_cols, inplace=True, errors="ignore")
65     X_val_df.drop(columns=const_cols, inplace=True, errors="ignore")
66     X_tr_df.drop(columns=const_cols, inplace=True, errors="ignore")
67     X_ts_df.drop(columns=const_cols, inplace=True, errors="ignore")
68
69 # 6) Scale with RobustScaler on benign sample
70 from sklearn.preprocessing import RobustScaler
71 rs = np.random.RandomState(42)
72 N_TRAIN_MAX = 400_000
73 idx_fit = rs.choice(len(X_train_df), size=min(N_TRAIN_MAX, len(X_train_df)), replace=False)
74 scaler = RobustScaler(with_centering=True, with_scaling=True)
75 scaler.fit(X_train_df.iloc[idx_fit].values)
76
77 def transform_df_to_array(Xdf):
78     X = scaler.transform(Xdf.values).astype("float32", copy=False)
79     # enforce finiteness just in case
80     X = np.nan_to_num(X, nan=0.0, posinf=1e6, neginf=-1e6)
81     return X
82
83 X_train = transform_df_to_array(X_train_df); del X_train_df
84 X_val   = transform_df_to_array(X_val_df);   del X_val_df
85 X_tr    = transform_df_to_array(X_tr_df);    del X_tr_df
86 X_ts    = transform_df_to_array(X_ts_df);    del X_ts_df
87 gc.collect()
88
89 # 7) Re-check finiteness
90 def check_finiteness(X, name):
91     print(name, "NaN:", np.isnan(X).sum(), "Inf:", np.isinf(X).sum(), "shape:", X.shape)
92     check_finiteness(X_train, "X_train")
93     check_finiteness(X_val,   "X_val")
94     check_finiteness(X_tr,   "X_tr")
95     check_finiteness(X_ts,   "X_ts")
96
97 # 8) If you want to train on a subset of X_train (benign), pick it now
98 N_FIT_MAX = 400_000
99 if len(X_train) > N_FIT_MAX:
100     idx_fit2 = rs.choice(len(X_train), size=N_FIT_MAX, replace=False)
101     X_fit = X_train[idx_fit2]
102 else:
103     X_fit = X_train
104 print("Fitting AE on:", X_fit.shape)
105

```

```

Dropping constant/invalid cols: ['is_sm_ips_ports']
X_train NaN: 0 Inf: 0 shape: (2218760, 40)
X_val NaN: 0 Inf: 0 shape: (761893, 40)
X_tr NaN: 0 Inf: 0 shape: (2219157, 40)
X_ts NaN: 0 Inf: 0 shape: (2219157, 40)
Fitting AE on: (400000, 40)

```

```

1 import tensorflow as tf
2 from tensorflow import keras
3 tf.random.set_seed(42)
4
5 input_dim = X_fit.shape[1]
6 inp = keras.Input(shape=(input_dim,))
7 x = keras.layers.Dense(256, activation='relu')(inp)
8 x = keras.layers.Dense(128, activation='relu')(x)
9 z = keras.layers.Dense(64, activation='relu')(x)
10 x = keras.layers.Dense(128, activation='relu')(z)
11 x = keras.layers.Dense(256, activation='relu')(x)
12 out = keras.layers.Dense(input_dim, activation='linear')(x)
13
14 ae = keras.Model(inp, out)
15 ae.compile(optimizer=keras.optimizers.Adam(learning_rate=5e-4, clipnorm=1.0), loss='mse')
16 es = keras.callbacks.EarlyStopping(monitor='loss', patience=5, restore_best_weights=True)
17 _ = ae.fit(X_fit, X_fit, epochs=60, batch_size=1024, shuffle=True, callbacks=[es], verbose=0)
18
19 # Reconstructions
20 val_rec = ae.predict(X_val, batch_size=4096, verbose=0)
21 tr_rec = ae.predict(X_tr, batch_size=4096, verbose=0)
22 ts_rec = ae.predict(X_ts, batch_size=4096, verbose=0)
23
24 # Errors (now guaranteed finite)
25 val_err = np.mean((X_val - val_rec)**2, axis=1).astype("float32")
26 tr_err = np.mean((X_tr - tr_rec)**2, axis=1).astype("float32")
27 ts_err = np.mean((X_ts - ts_rec)**2, axis=1).astype("float32")
28
29 # Threshold selection & evaluation (same functions you already have)
30 th_f2_ae, p_f2_ae, r_f2_ae, f2_ae = pick_thresh_by_Fbeta(y_val, val_err, beta=2.0)
31 y_tr_f2_ae = (tr_err >= th_f2_ae).astype(int)
32 y_ts_f2_ae = (ts_err >= th_f2_ae).astype(int)
33 print(f"AE F2 threshold: {th_f2_ae:.6f} (val P={p_f2_ae:.3f}, R={r_f2_ae:.3f}, F2={f2_ae:.3f})")
34 summarize(y_tr, y_tr_f2_ae, tr_err, "AE | REALISTIC | F2-opt")
35 summarize(y_ts, y_ts_f2_ae, ts_err, "AE | STRESS | F2-opt")
36
37 th_fpr_ae, fpr_sel_ae, tpr_sel_ae = pick_thresh_at_fpr(y_val, val_err, max_fpr=0.05)
38 y_tr_fpr_ae = (tr_err >= th_fpr_ae).astype(int)
39 y_ts_fpr_ae = (ts_err >= th_fpr_ae).astype(int)
40 print(f"AE FPR<=5% threshold: {th_fpr_ae:.6f} (val FPR={fpr_sel_ae:.3f}, TPR={tpr_sel_ae:.3f})")
41 summarize(y_tr, y_tr_fpr_ae, tr_err, "AE | REALISTIC | FPR<=5%")
42 summarize(y_ts, y_ts_fpr_ae, ts_err, "AE | STRESS | FPR<=5%")

```

⁴³ Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.